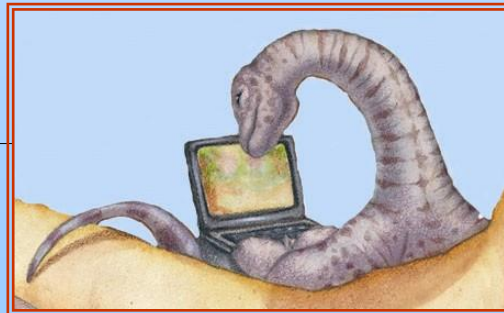


The picture can't be displayed.

Chapter 3: Processes






Process Concept

- ❑ An operating system executes a variety of programs:
 - ❑ Batch system – jobs
 - ❑ Time-shared systems – user programs or tasks
- ❑ Textbook uses the terms *job* and *process* almost interchangeably
- ❑ Process – a program in execution; process execution must progress in sequential fashion
- ❑ A process includes:
 - ❑ program counter
 - ❑ stack
 - ❑ data section





Process in Memory





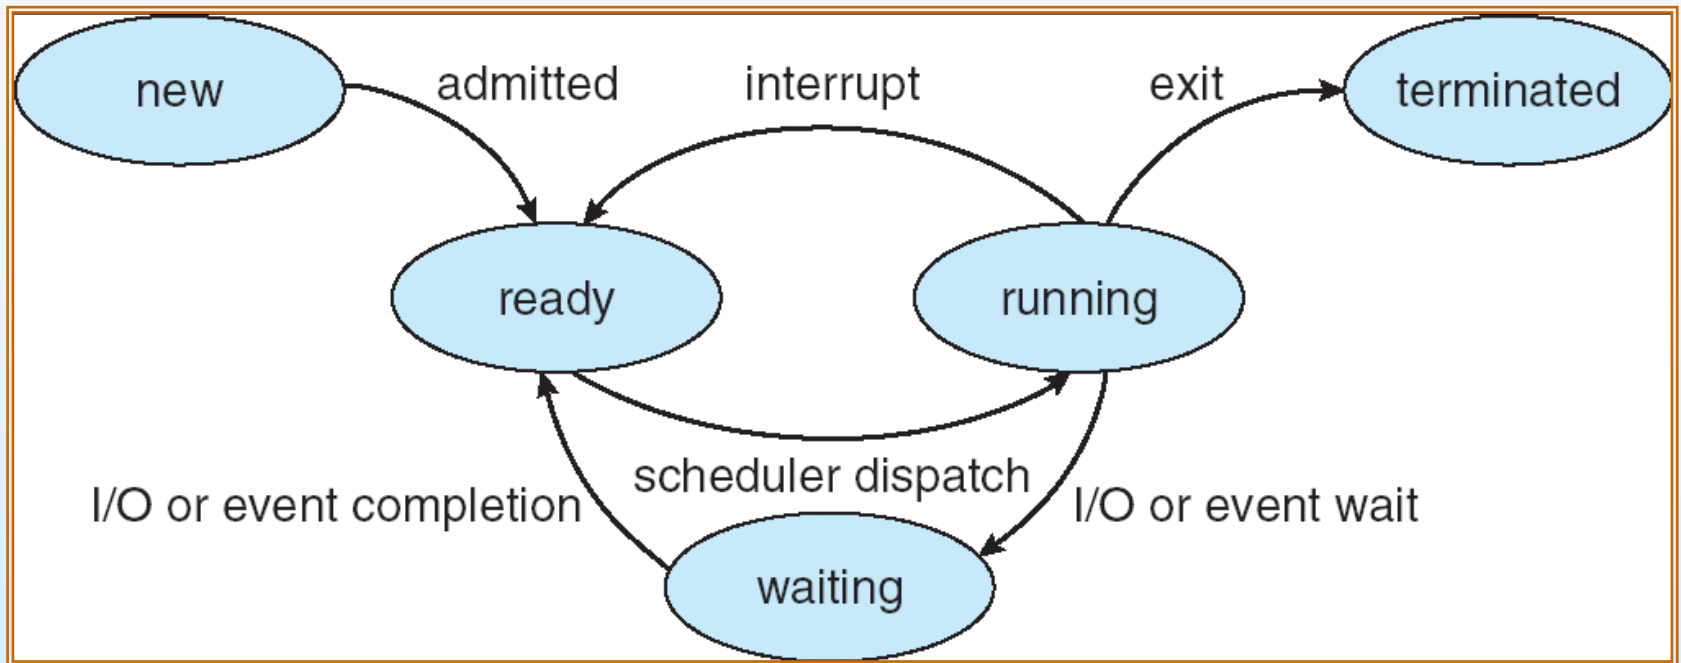
Process State

- As a process executes, it changes *state*
 - **new**: The process is being created
 - **running**: Instructions are being executed
 - **waiting**: The process is waiting for some event to occur
 - **ready**: The process is waiting to be assigned to a process
 - **terminated**: The process has finished execution





Diagram of Process State





Process Control Block (PCB)

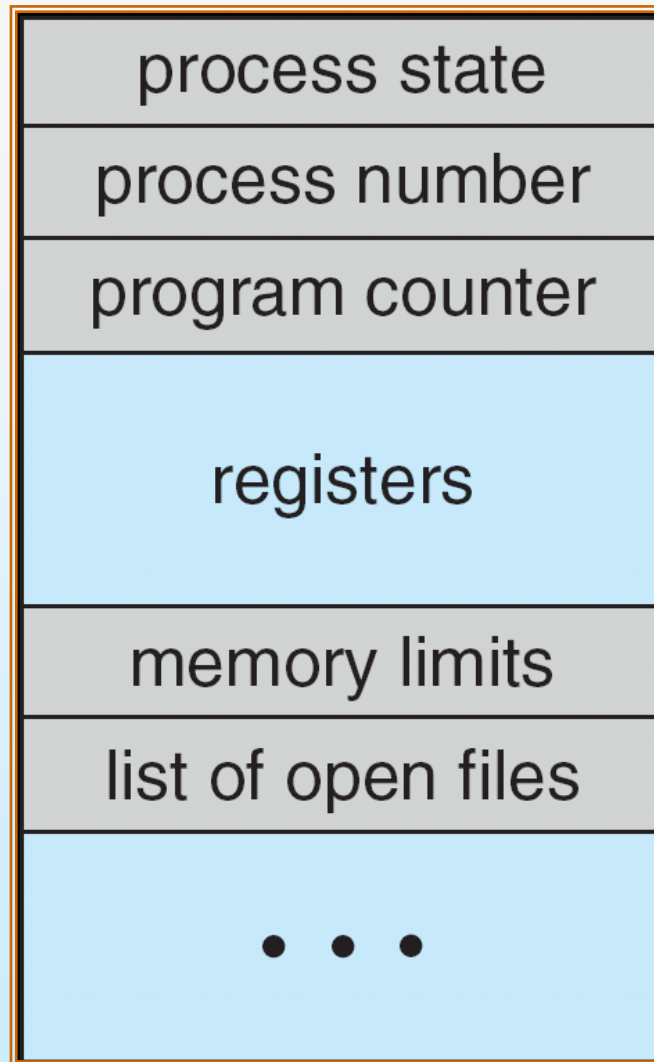
Information associated with each process

- Process state
- Program counter
- CPU registers
- CPU scheduling information
- Memory-management information
- Accounting information
- I/O status information



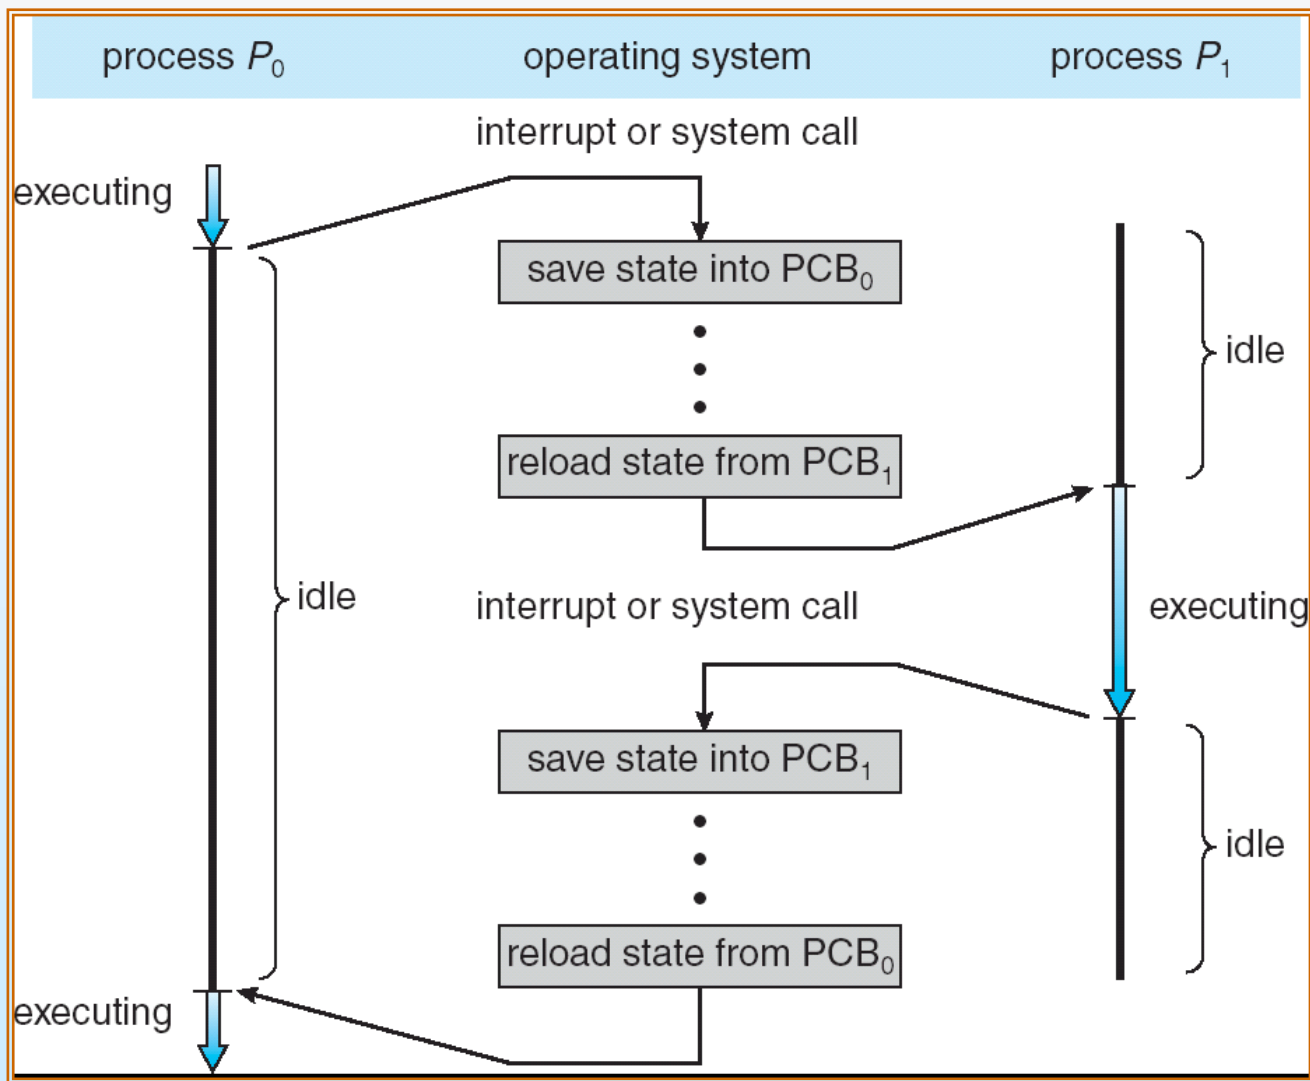


Process Control Block (PCB)





CPU Switch From Process to Process





Process Scheduling Queues

- ❑ **Job queue** – set of all processes in the system
- ❑ **Ready queue** – set of all processes residing in main memory, ready and waiting to execute
- ❑ **Device queues** – set of processes waiting for an I/O device
- ❑ Processes migrate among the various queues



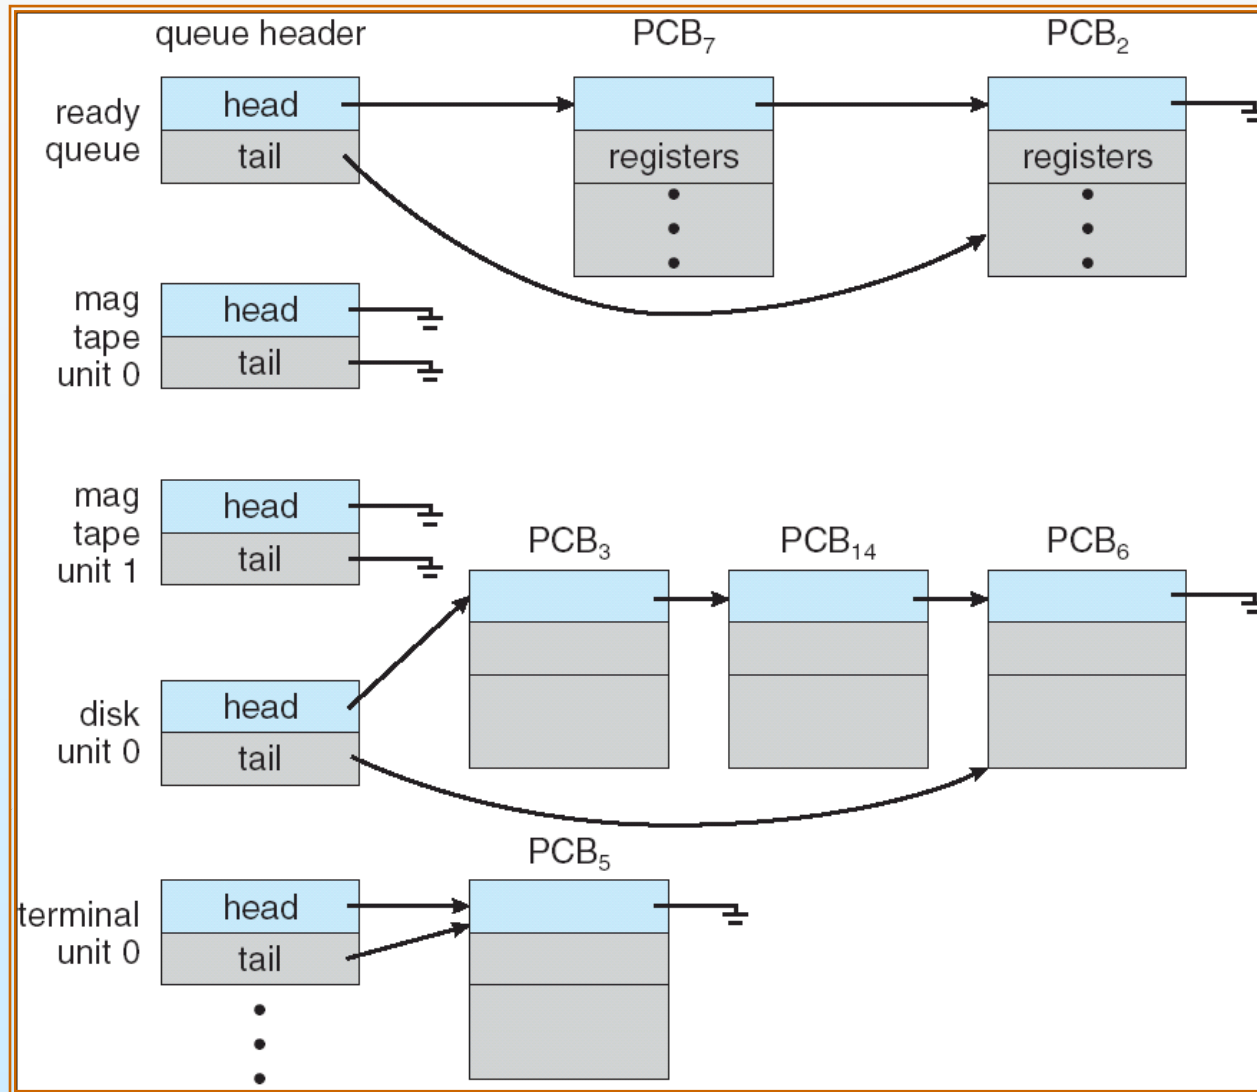


Chap03.pdf



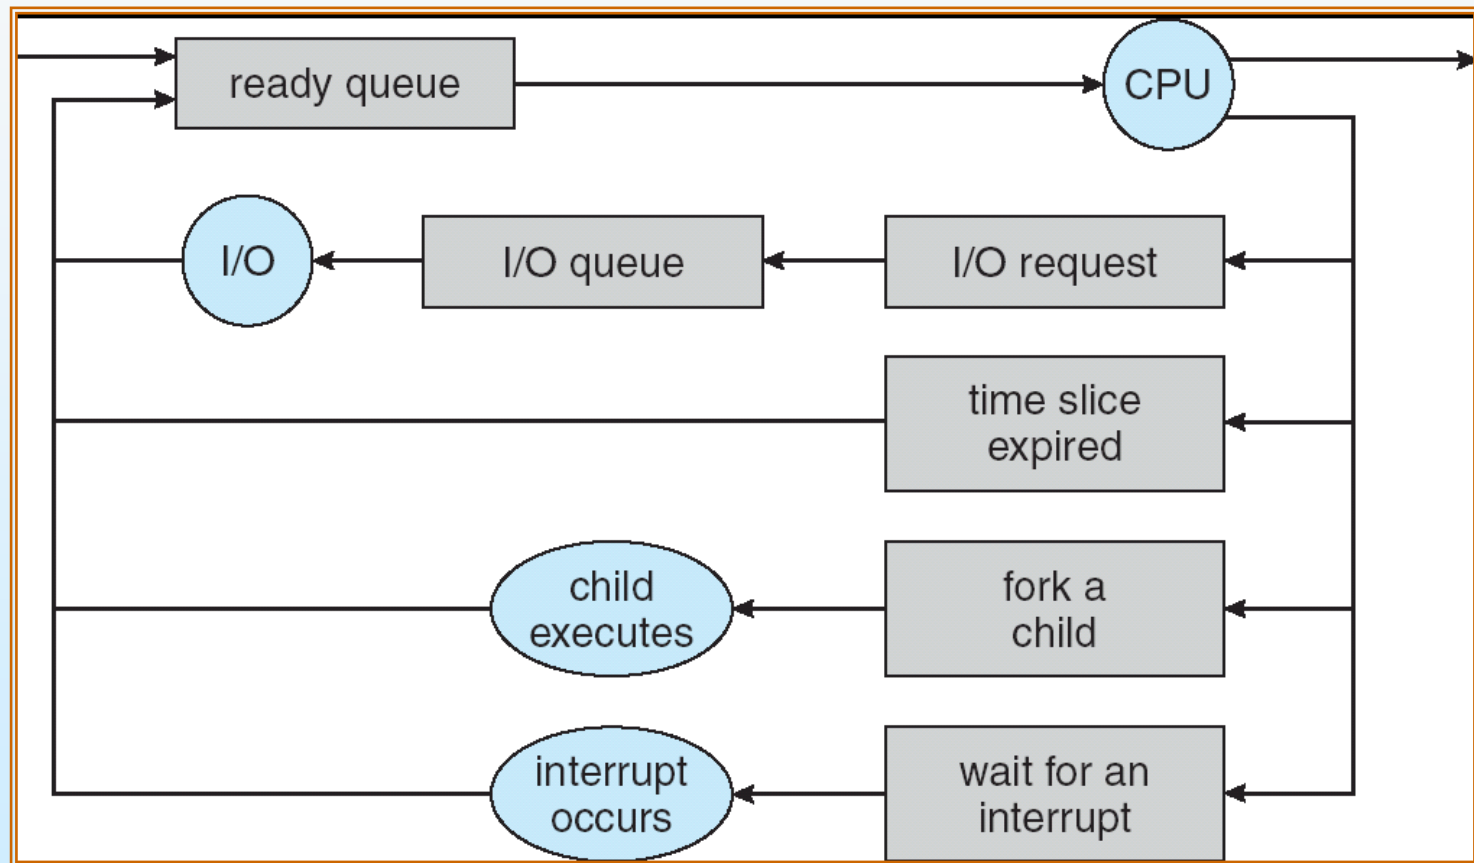


Ready Queue And Various I/O Device Queues





Representation of Process Scheduling





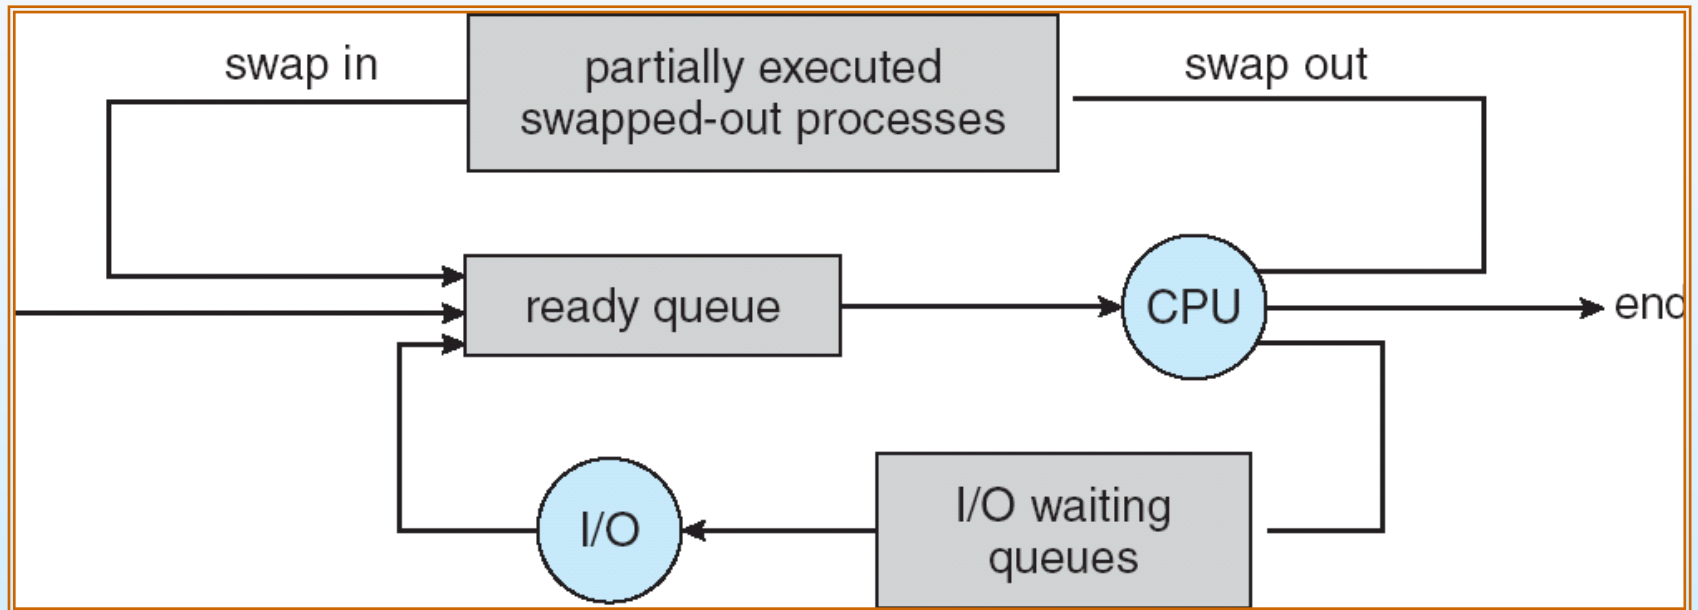
Schedulers

- **Long-term scheduler** (or job scheduler) – selects which processes should be brought into the ready queue
- **Short-term scheduler** (or CPU scheduler) – selects which process should be executed next and allocates CPU





Addition of Medium Term Scheduling





Schedulers (Cont.)

- Short-term scheduler invoked very frequently (milliseconds) \Rightarrow (must be fast)
- Long-term scheduler invoked very infrequently (seconds, minutes) \Rightarrow (may be slow)
- Long-term scheduler controls *degree of multiprogramming*
- Mix of:
 - **I/O-bound process** – spends more time doing I/O than computations, many short CPU bursts
 - **CPU-bound process** – spends more time doing computations; few very long CPU bursts





Context Switch

- CPU switches to another process:
 - must save the state of the old process
 - load the saved state for the new process

- Context-switch time: **overhead**;
 - no useful work while switching

- Time dependent on hardware support





Process Creation

- Parent process create children processes, which, in turn create other processes (**tree** of processes)
- Resource sharing **POLICIES**
 - Parent/children **share all** resources
 - Children **share subset** of parent's resources
 - Parent and child **share no** resources
- Execution **POLICIES**
 - Parent and children execute **concurrently**
 - Parent **waits** until children terminate





Process Creation (Cont.)

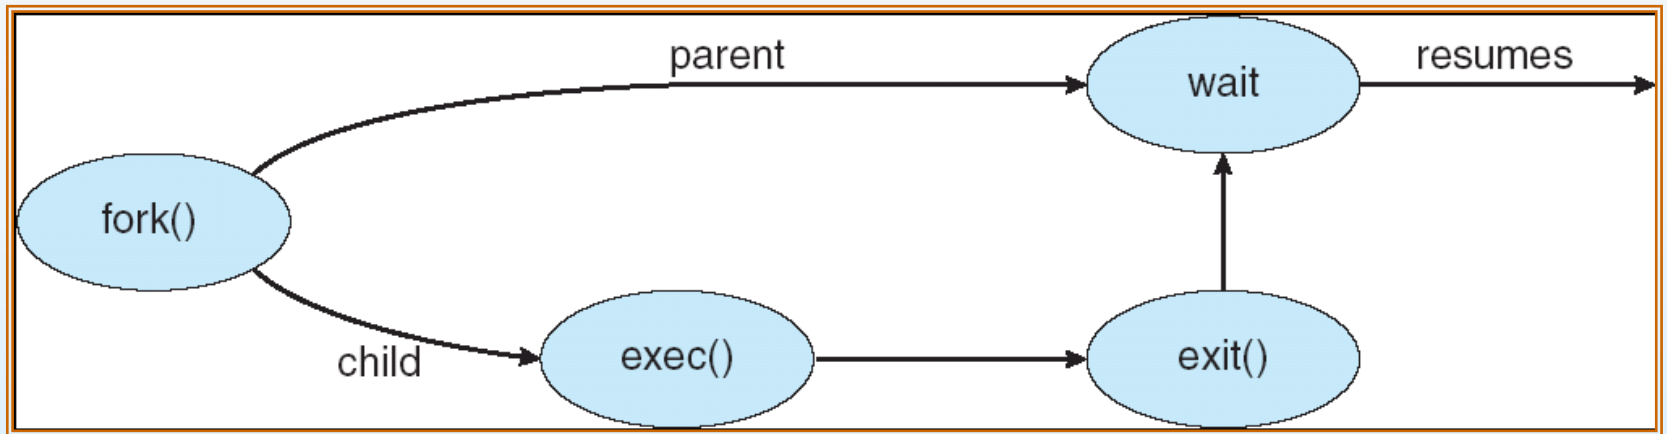
- Address space **POLICIES**
 - Child **duplicate** of parent
 - Child has a **new program** loaded into it

- UNIX examples
 - **fork** system call creates new process
 - **exec** system call used (in general, after a **fork**)
to replace process' memory space with a new program





Process Creation





C Program Forking Separate Process

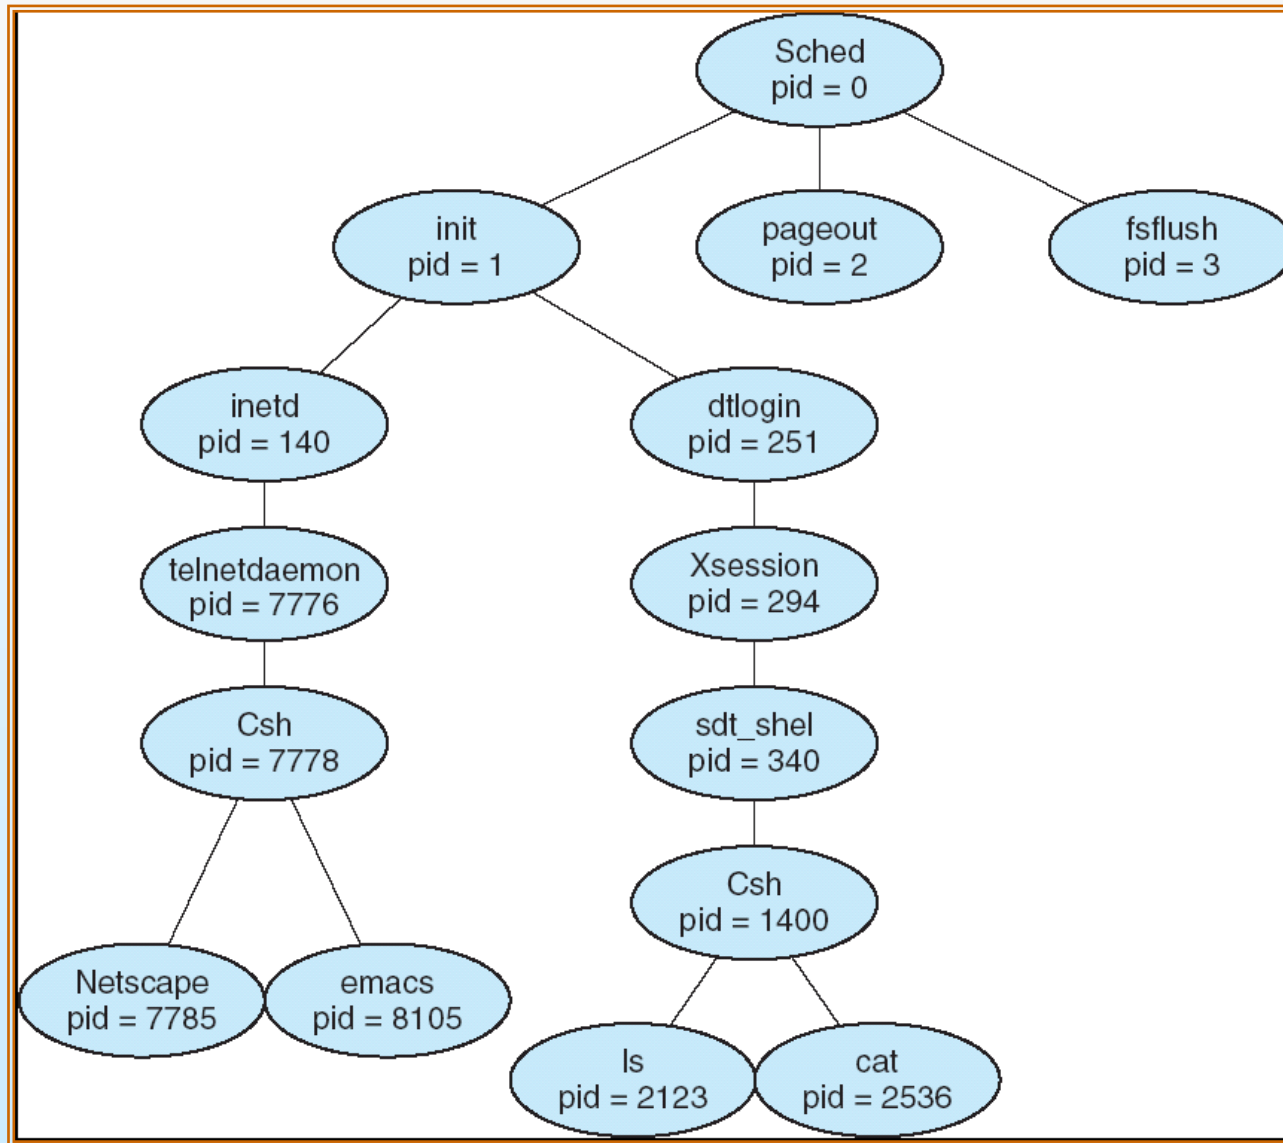
```
int main() {
    Pid_t  pid;

    pid = fork();      /* fork another process */
    if (pid == 0) { /* child process */
        execlp("/bin/ls", "ls", "-la", NULL);
    }
    else if (pid > 0) { /* parent */
        wait (NULL);  /* wait for child to complete */
        printf ("Child Complete");
        exit(0);
    }
    else if (pid < 0) { /* error occurred */
        fprintf(stderr, "Fork Failed");
        exit(-1);
    }
}
```





A tree of processes on a typical Solaris





Process Termination

- Process executes last statement and asks the operating system to delete it (**exit**)
 - Output data from child to parent (via **wait**)
 - Process' resources are deallocated by operating system
- Parent may terminate execution of children processes (**abort**)
 - Child has exceeded allocated resources
 - Task assigned to child is no longer required
 - If parent is exiting
 - ▶ Some operating system do not allow child to continue if its parent terminates
 - All children terminated - *cascading termination*





Interprocess Communication (IPC)

- ❑ Mechanism for processes to communicate and synchronize their actions
- ❑ Message passing – processes communicate with each other without resorting to shared variables
- ❑ IPC facility provides two operations:
 - ❑ **send**(*message*) – message size fixed or variable
 - ❑ **receive**(*message*)
- ❑ If P and Q wish to communicate, they need to:
 - ❑ establish a *communication link* between them
 - ❑ exchange messages via send/receive
- ❑ Implementation of communication link
 - ❑ physical (e.g., shared memory, hardware bus)
 - ❑ logical (e.g., logical properties)





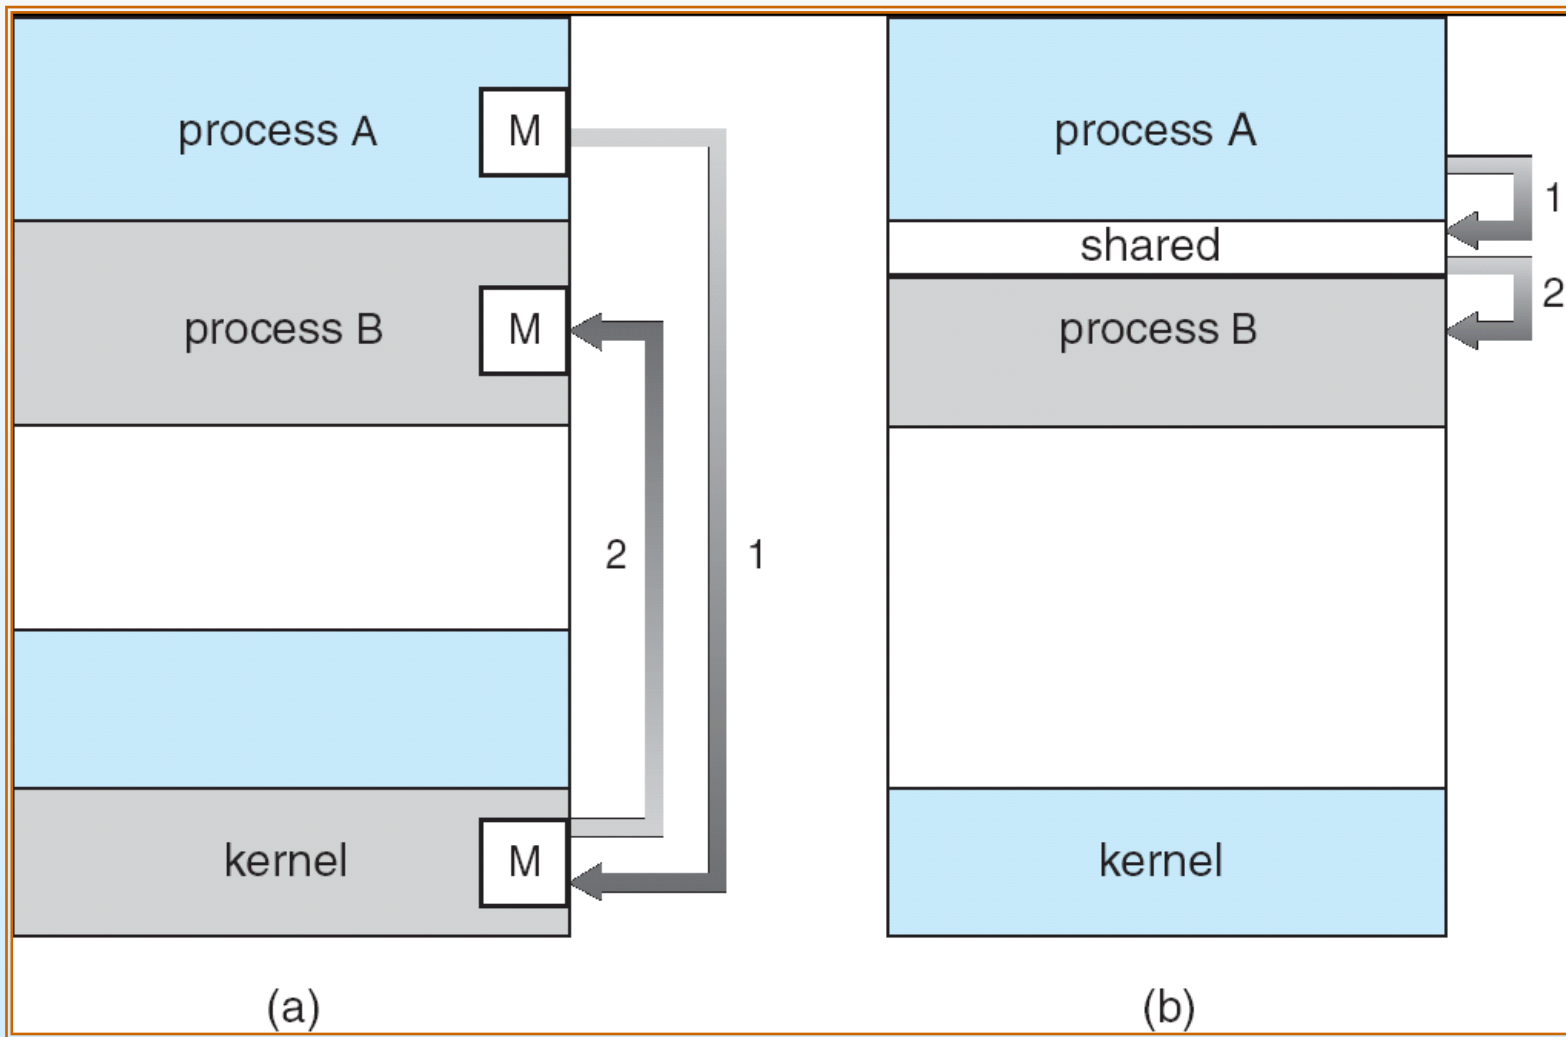
Implementation Questions

- How are links established?
- Can a link be associated with more than two processes?
- How many links can there be between every pair of communicating processes?
- What is the capacity of a link?
- Is the size of a message for the fixed or variable?
- Is a link unidirectional or bi-directional?





Communications Models





Direct Communication

- Processes must name each other explicitly:
 - **send** (P , $message$) – send a message to process P
 - **receive**(Q , $message$) – receive a message from process Q

- Properties of communication link
 - Links established automatically
 - A link is associated with exactly one pair of communicating processes
 - Between each pair there exists exactly one link
 - The link may be unidirectional, but is usually bi-directional





Indirect Communication

- Messages are directed/received from **mailboxes** (also referred to as **ports**)
 - Each mailbox has a unique id
 - Processes can communicate only if they share a mailbox

- Properties of communication link
 - Link established only if processes share a common mailbox
 - A link may be associated with many processes
 - Each pair of processes may share several communication links
 - Link may be unidirectional or bi-directional





Indirect Communication

- Operations
 - create a new mailbox
 - send and receive messages through mailbox
 - destroy a mailbox

- Primitives are defined as:
 - send**(*A, message*) – send a message to mailbox *A*
 - receive**(*A, message*) – receive a message from mailbox *A*





Indirect Communication

- Mailbox sharing
 - P_1 , P_2 , and P_3 share mailbox A
 - P_1 , sends; P_2 and P_3 receive
 - Who gets the message?

- Solutions
 - Allow a link to be associated with at most two processes
 - Allow only one process at a time to execute a receive operation
 - Allow the system to select arbitrarily the receiver. Sender is notified who the receiver was.





Synchronization

- Message passing may be either blocking or non-blocking
- **Blocking** is considered **synchronous**
 - **Blocking send** has the sender block until the message is received
 - **Blocking receive** has the receiver block until a message is available
- **Non-blocking** is considered **asynchronous**
 - **Non-blocking send** has the sender send the message and continue
 - **Non-blocking receive** has the receiver receive a valid message or null





Buffering

- Queue of messages attached to the link; implemented in one of three ways
 1. Zero capacity – 0 messages
Sender must wait for receiver (rendezvous)
 2. Bounded capacity – finite length of n messages
Sender must wait if link full
 3. Unbounded capacity – infinite length
Sender never waits





Client-Server Communication

- Sockets
- Remote Procedure Calls
- Remote Method Invocation (Java)





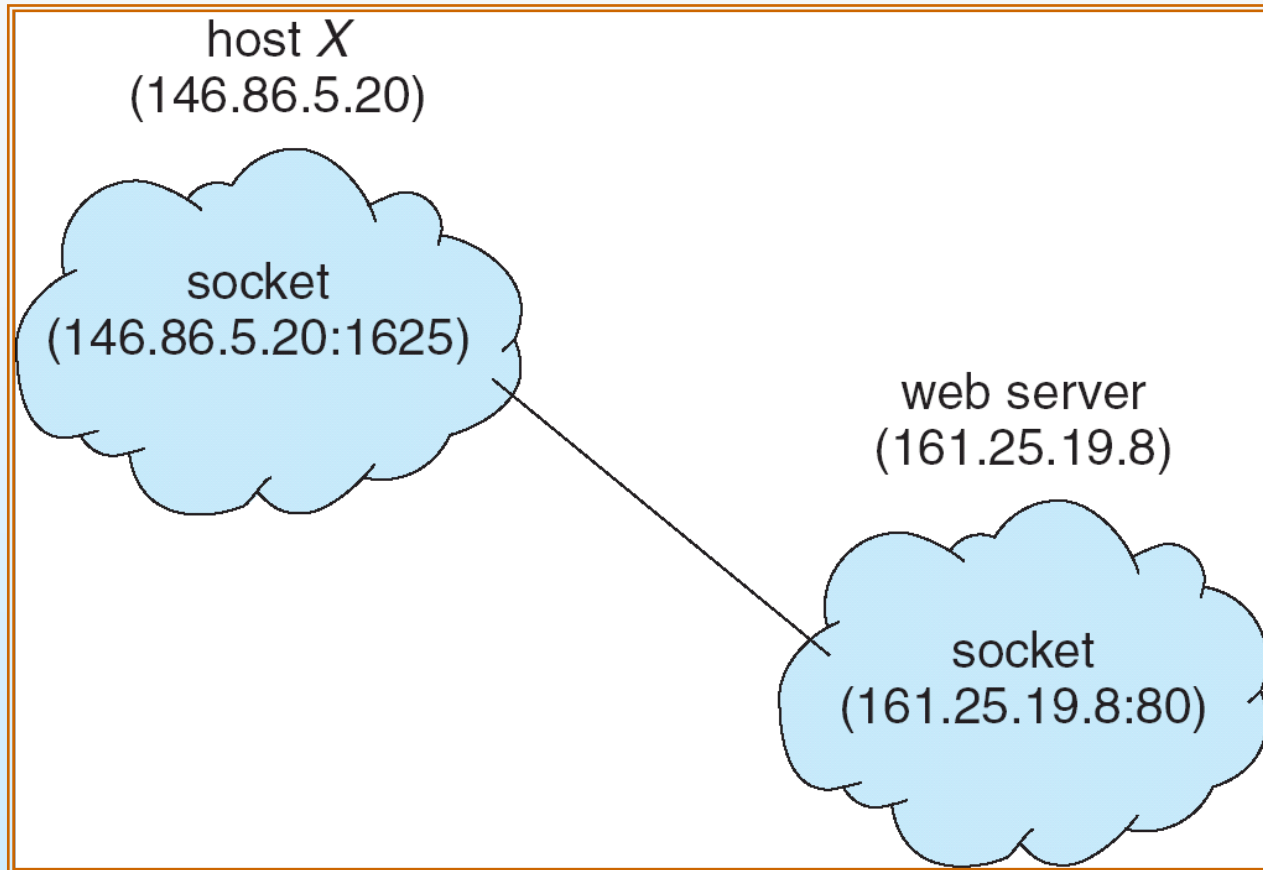
Sockets

- A socket is defined as an *endpoint for communication*
- Concatenation of IP address and port
- The socket **161.25.19.8:1625** refers to port **1625** on host **161.25.19.8**
- Communication consists between a pair of sockets





Socket Communication





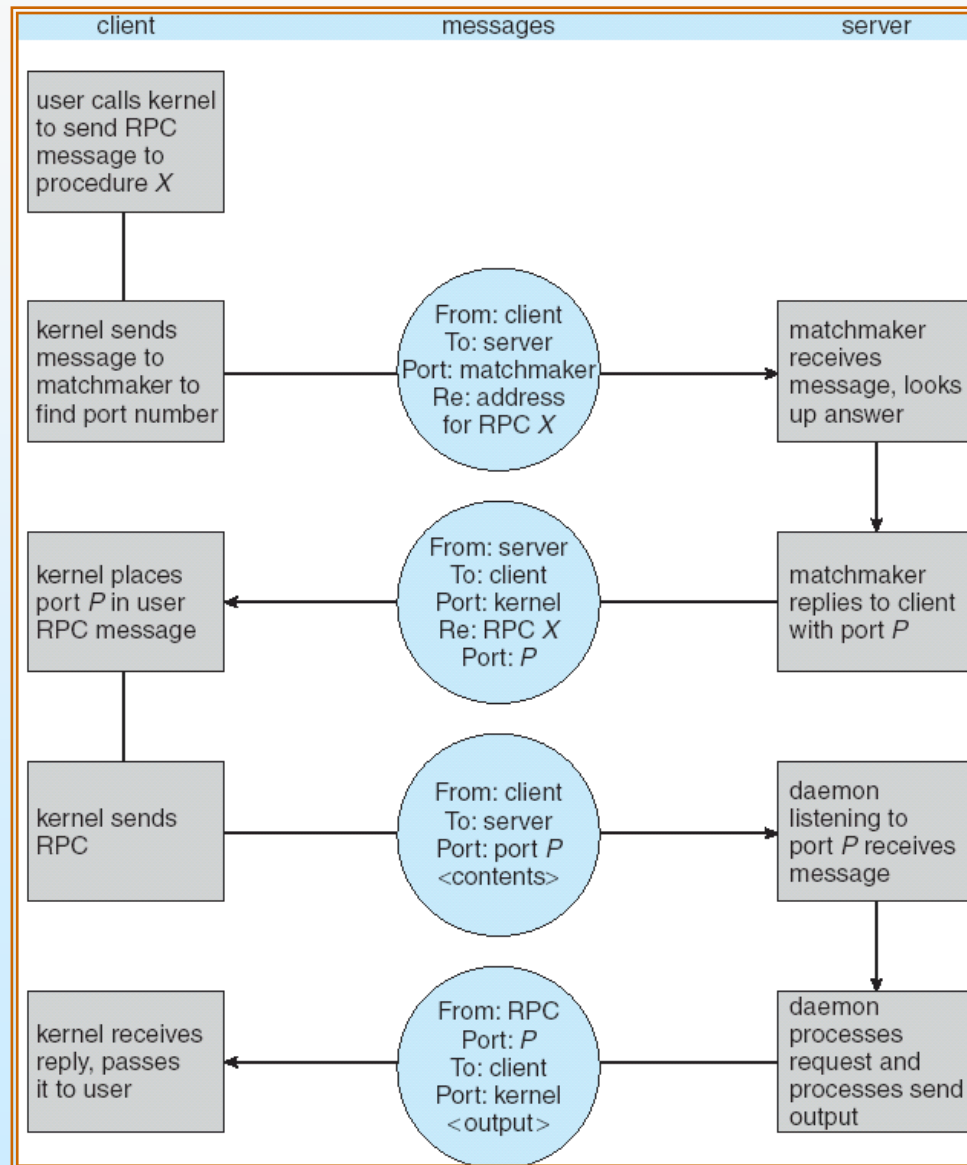
Remote Procedure Calls (RPC)

- RPC abstracts procedure calls between processes on networked systems.
- **Stubs** – client-side proxy for the actual procedure on the server.
- The client-side stub locates the server and *marshalls* the parameters.
- Server-side stub receives this message, unpacks the marshalled parameters, and performs the procedure on the server.






Execution of RPC





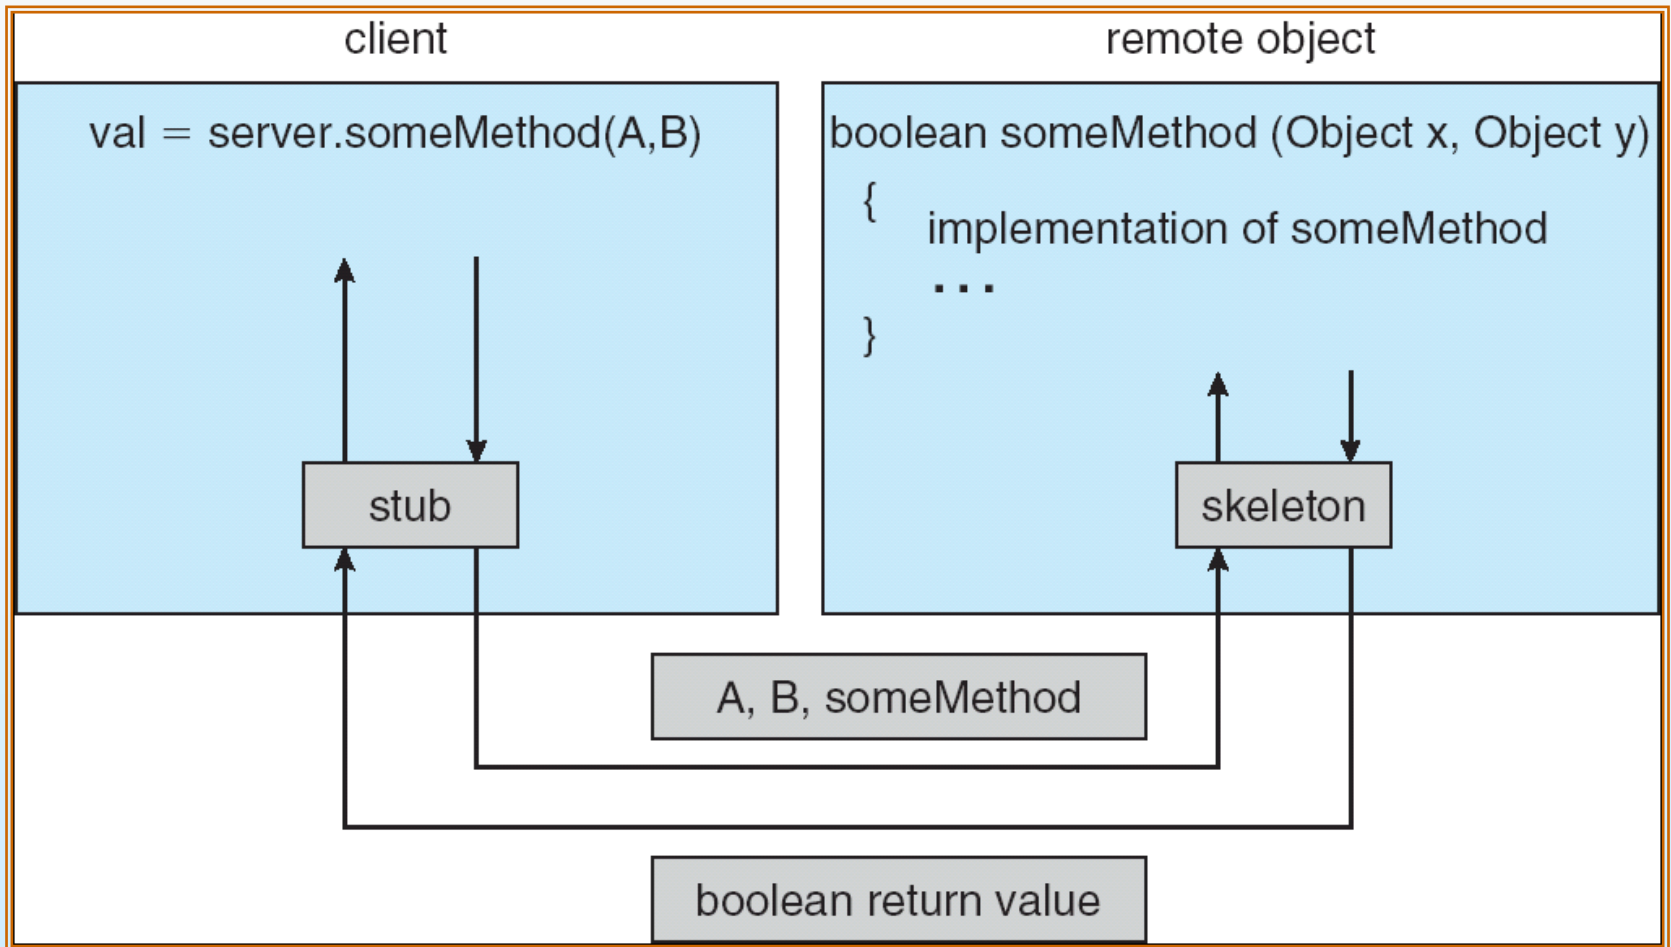
Remote Method Invocation (RMI)

- ❑ RMI is a Java mechanism similar to RPCs.
- ❑ RMI allows a Java program on one machine to invoke a method on a remote object.





Marshalling Parameters



The picture can't be displayed.

End of Chapter 3

