

Process Description and Control

Chapter 3

All multiprogramming OS are build around the concept of processes

A process is sometimes called a task

OS Requirements for Processes

- **OS must interleave the execution of several processes to maximize CPU usage while providing reasonable response time**
- **OS must allocate resources to processes while avoiding deadlock**
- **OS must support inter process communication and user creation of processes**

Dispatcher (short-term scheduler)

- Is an OS program that moves the processor from one process to another
- It prevents a single process from monopolizing processor time
- It decides who goes next according to a scheduling algorithm (to be discussed later)
- The CPU will always execute instructions from the dispatcher while switching from process A to process B

43

When does a process gets created?

- Submission of a batch job
- User logs on
- Created by OS to provide a service to a user (ex: printing a file)
- Spawned by an existing process
 - ◆ a user program can dictate the creation of a number of processes

44

When does a process gets terminated?

- Batch job issues *Halt* instruction
- User logs off
- Process executes a service request to terminate
- Error and fault conditions

45

Reasons for Process Termination

- Normal completion
- Time limit exceeded
- Memory unavailable
- Memory bounds violation
- Protection error
 - ◆ example: write to read-only file
- Arithmetic error
- Time overrun
 - ◆ process waited longer than a specified maximum for an event

46

Reasons for Process Termination

- **I/O failure**
- **Invalid instruction**
 - ◆ happens when try to execute data
- **Privileged instruction**
- **Operating system intervention**
 - ◆ such as when deadlock occurs
- **Parent request to terminate one offspring**
- **Parent terminates so child processes terminate**

47

Process States

- **Let us start with these states:**
 - ◆ The Running state
 - ☞ The process that gets executed (single CPU)
 - ◆ The Ready state
 - ☞ any process that is ready to be executed
 - ◆ The Blocked state
 - ☞ when a process cannot execute until some event occurs (ex: the completion of an I/O)

48

Other Useful States

■ The New state

- ◆ OS has performed the necessary actions to create the process
 - ☞ has created a process identifier
 - ☞ has created tables needed to manage the process
- ◆ but has not yet committed to execute the process (not yet admitted)
 - ☞ because resources are limited

49

Other Useful States

■ The Exit state

- ◆ Termination moves the process to this state
- ◆ It is no longer eligible for execution
- ◆ Tables and other info are temporarily preserved for auxiliary program
 - ☞ Ex: accounting program that cumulates resource usage for billing the users
- **The process (and its tables) gets deleted when the data is no more needed**

50

Process Transitions

- **Ready --> Running**

- ◆ When it is time, the dispatcher selects a new process to run

- **Running --> Ready**

- ◆ the running process has expired his time slot
- ◆ the running process gets interrupted because a higher priority process is in the ready state

51

Process Transitions

- **Running --> Blocked**


- ◆ When a process requests something for which it must wait
 - ☞ a service that the OS is not ready to perform
 - ☞ an access to a resource not yet available
 - ☞ initiates I/O and must wait for the result
 - ☞ waiting for a process to provide input (IPC)

- **Blocked --> Ready**

- ◆ When the event for which it was waiting occurs

52

A Five-state Process Model



Ready to exit: A parent may terminate a child process

53

The need for swapping

- So far, all the processes had to be (at least partly) in main memory
- Even with virtual memory, keeping too many processes in main memory will deteriorate the system's performance
- The OS may need to suspend some processes, ie: to swap them out to disk. We add 2 new states:
- Blocked Suspend: blocked processes which have been swapped out to disk
- Ready Suspend: ready processes which have been swapped out to disk

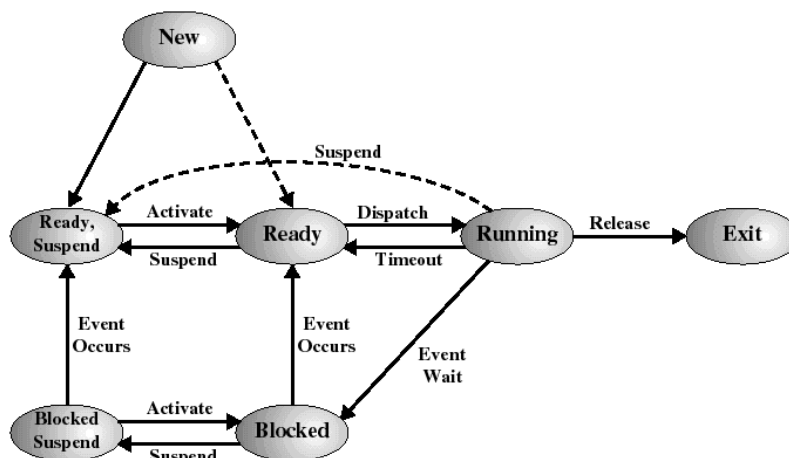
54

New state transitions (mid-term scheduling)

- **Blocked --> Blocked Suspend**
 - ◆ When all processes are blocked, the OS will make room to bring a ready process in memory
- **Blocked Suspend --> Ready Suspend**
 - ◆ When the event for which it has been waiting occurs (state info is available to OS)
- **Ready Suspend --> Ready**
 - ◆ when no more ready process in main memory
- **Ready --> Ready Suspend (unlikely)**
 - ◆ When there are no blocked processes and must free memory for adequate performance

55

A Seven-state Process Model



56

Operating System Control Structures

- **An OS maintains the following tables for managing processes and resources:**
 - ◆ Memory tables (to be discussed later)
 - ◆ I/O tables (to be discussed later)
 - ◆ File tables (to be discussed later)
 - ◆ Process tables

57

Process Image

- **User program**
- **User data**
- **Stack(s)**
 - ◆ for procedure calls and parameter passing
- **Process Control Block (execution context)**
 - ◆ Data needed (process attributes) by the OS to control the process. This includes:
 - ☞ Process identification information
 - ☞ Processor state information
 - ☞ Process control information

58

Location of the Process Image

- **Each process image is in virtual memory**
 - ◆ may not occupy a contiguous range of addresses (depends on the memory management scheme used)
 - ◆ both a private and shared memory address space is used
- **The location of each process image is pointed to by an entry in the Primary Process Table**
- **For the OS to manage the process, at least part of its image must be loaded into main memory**

59

Process Identification (in the PCB)

- **A few numeric identifiers may be used**
 - ◆ Unique process identifier (always)
 - ☞ indexes (directly or indirectly) into the primary process table
 - ◆ User identifier
 - ☞ the user who is responsible for the job
 - ◆ Identifier of the process that created this process

60

Processor State Information (in PCB)

- **Contents of processor registers**
 - ◆ User-visible registers
 - ◆ Control and status registers
 - ◆ Stack pointers
- **Program status word (PSW)**
 - ◆ contains status information
 - ◆ Example: the EFLAGS register on Pentium machines

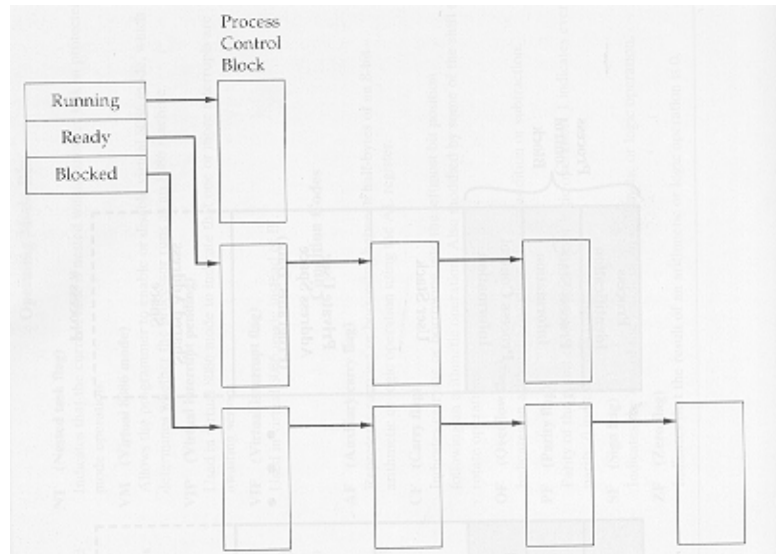
61

Process Control Information (in PCB)

- **scheduling and state information**
 - ◆ Process state (ie: running, ready, blocked...)
 - ◆ Priority of the process
 - ◆ Event for which the process is waiting (if blocked)
- **data structuring information**
 - ◆ may hold pointers to other PCBs for process queues, parent-child relationships and other structures

62

Queues as linked lists of PCBs



63

Process Control Information (in PCB)

- **interprocess communication**
 - ◆ may hold flags and signals for IPC
- **process privileges**
 - ◆ Ex: access to certain memory locations...
- **memory management**
 - ◆ pointers to segment/page tables assigned to this process
- **resource ownership and utilization**
 - ◆ resource in use: open files, I/O devices...
 - ◆ history of usage (of CPU time, I/O...)

64

Modes of Execution

- **To provide protection to PCBs (and other OS data) most processors support at least 2 execution modes:**
 - ◆ Privileged mode (a.k.a. system mode, kernel mode, supervisor mode, control mode)
 - ☞ manipulating control registers, primitive I/O instructions, memory management...
 - ◆ User mode
- **For this the CPU provides a (or a few) mode bit which may only be set by an interrupt or trap or OS call**

65

Process Creation

- **Assign a unique process identifier**
- **Allocate space for the process image**
- **Initialize process control block**
 - ◆ many default values (ex: state is New, no I/O devices or files...)
- **Set up appropriate linkages**
 - ◆ Ex: add new process to linked list used for the scheduling queue

66

When to Switch a Process ?

- **A process switch may occur whenever the OS has gained control of CPU. ie when:**
 - ◆ **Supervisor Call**
 - ☞ explicit request by the program (ex: file open).
The process will probably be blocked
 - ◆ **Trap**
 - ☞ An error resulted from the last instruction. It may cause the process to be moved to the Exit state
 - ◆ **Interrupt**
 - ☞ the cause is external to the execution of the current instruction. Control is transferred to IH

67

Mode Switching

- **It may happen that an interrupt does not produce a process switch**
- **The control can just return to the interrupted program**
- **Then only the processor state information needs to be saved on stack**
- **This is called mode switching (user to kernel mode when going into IH)**
- **Less overhead: no need to update the PCB like for process switching**

68

Steps in Process (Context) Switching

- **Save context of processor including program counter and other registers**
- **Update the PCB of the running process with its new state and other associate info**
- **Move PCB to appropriate queue - ready, blocked**
- **Select another process for execution**
- **Update PCB of the selected process**
- **Restore CPU context from that of the selected process**


69

Execution of the Operating System

- **Up to now, by process we were referring to “user process”**
- **If the OS is just like any other collection of programs, is the OS a process?**
- **If so, how it is controlled?**
- **The answer depends on the OS design.**

70

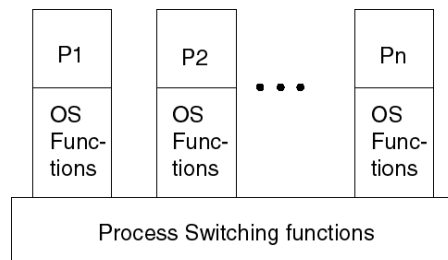
Nonprocess Kernel (old)



- The concept of process applies only to user programs
- OS code is executed as a separate entity in privilege mode
- OS code never gets executed within a process

71

Execution within User Processes

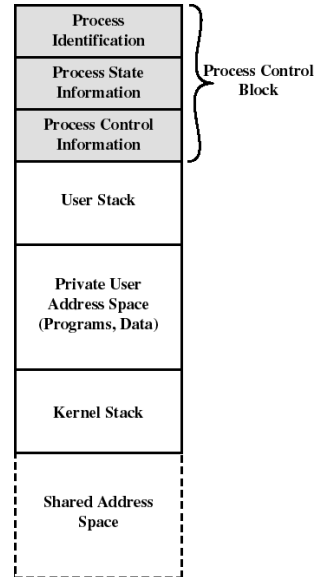


- Virtually all OS code gets executed within the context of a user process
- On Interrupts, Traps, System calls: the CPU switch to kernel mode to execute OS routine within the context of user process (mode switch)
- Control passes to process switching functions (outside processes) only when needed

72


Execution within User Processes

- OS code and data are in the shared address space and are shared by all user processes
- Separate kernel stack for calls/returns when the process is in kernel mode
- Within a user process, both user and OS programs may execute (more than 1)



73

Process-based Operating System



- The OS is a collection of system processes
- major kernel functions are separate processes
- small amount of process switching functions is executed outside of any process
- Design that easily makes use of multiprocessors

74