

# Memory Management

Chapter 7

# Memory Management

---

- ▶ OS + hardware: subdivide memory to accommodate multiple processes
- ▶ Memory: allocated to ensure reasonable supply of ready processes (consume available CPU time)
- ▶ Few processes in main memory?
  - ▶ Much of the time all processes will be waiting for I/O
  - ▶ CPU will be idle
- ▶ Memory needs to be allocated efficiently in order to pack as many processes into memory as possible

# Memory Management

---

- ▶ In most schemes, kernel occupies some fixed portion of main memory and rest shared by multiple processes

# Memory Management Requirements

## ► Relocation

- Programmer: not know where program will be in memory when executed
- process may be **relocated** in main memory due to swapping
- swapping enables the OS to have a larger pool of ready-to-execute processes
- memory references in code (for both instructions and data) must be translated to actual physical memory address

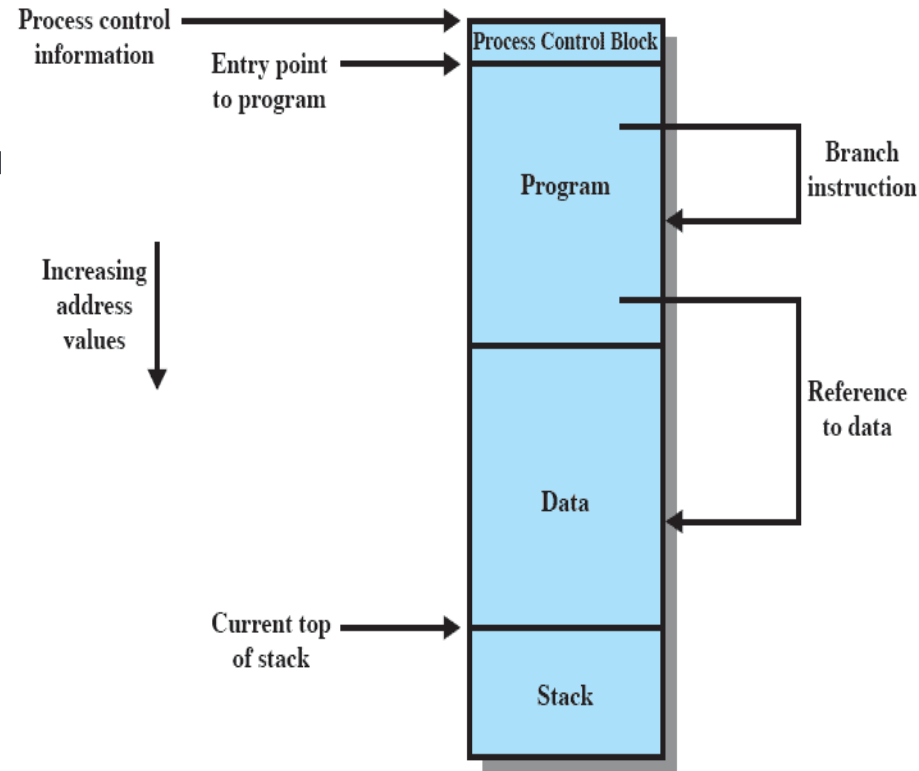


Figure 7.1 Addressing Requirements for a Process

# Memory Management Requirements

---

## ▶ Protection

- ▶ processes should not be able to reference memory locations in another process without permission
- ▶ impossible to check addresses at compile time in programs since the program could be relocated
- ▶ address references must be checked at run time by hardware
  - ▶ Operating system cannot anticipate all of the memory references a program will make



# Memory Management Requirements

---

## ▶ Sharing

- ▶ must allow several processes to access a common portion of main memory without compromising protection
  - ▶ cooperating processes may need to share access to the same data structure
  - ▶ better to allow each process to access the same copy of the program rather than have their own separate copy, e.g., parent and child processes after `fork()`.

# Memory Management Requirements

---

## ▶ Logical Organization

- ▶ Programs are written in modules
- ▶ Modules can be written and compiled independently
- ▶ Different degrees of protection given to modules (read-only, execute-only)
- ▶ Share modules among processes
- ▶ To effectively deal with user programs, the OS and hardware should support a basic form of module to provide the required protection and sharing

# Memory Management Requirements

---

## ▶ Physical Organization

- ▶ Secondary memory is the long term store for programs and data while main memory holds program and data currently in use
- ▶ Memory available for a program plus its data may be insufficient. Moving information between these two levels of memory is a major concern of memory management (OS)
  - ▶ it is highly inefficient to leave this responsibility to the application programmer
  - ▶ Overlaying allows various modules to be assigned the same region of memory
- ▶ Programmer does not know how much space will be available




# Simple Memory Management

---


- ▶ First we study the simpler case where there is **no virtual memory**
- ▶ An executing process must be loaded entirely in main memory (if overlays are not used)
- ▶ Although the following simple memory management techniques are not used in modern OS, they lay the ground for a proper discussion of virtual memory (to be discussed later)
  - ▶ fixed partitioning
  - ▶ dynamic partitioning
  - ▶ simple paging
  - ▶ simple segmentation

# Fixed Partitioning

- ▶ Partition main memory into a set of non overlapping regions called **partitions**
- ▶ Partitions can be of equal or unequal sizes



Equal-size partitions



Unequal-size partitions

# Fixed Partitioning

---

- ▶ any process whose size is less than or equal to a partition size can be loaded into the partition
- ▶ if all partitions are occupied, the operating system can swap a process out of a partition
- ▶ a program may be too large to fit in a partition. The programmer must then design the program with overlays
  - ▶ when the module needed is not present the user program must load that module into the program's partition, overlaying whatever program or data are there

# Fixed Partitioning

---

- ▶ Main memory use is inefficient. Any program, no matter how small, occupies an entire partition. This is called **internal fragmentation**.
- ▶ Unequal-size partitions lessens these problems but they still remain...
- ▶ Equal-size partitions was used in early IBM's OS/MFT (Multiprogramming with a Fixed number of Tasks)

# Placement Algorithm with Partitions


---

## ▶ Equal-size partitions

- ▶ If there is an available partition, a process can be loaded into that partition
  - ▶ because all partitions are of equal size, it does not matter which partition is used
- ▶ If all partitions are occupied by blocked processes, choose one process to swap out to make room for the new process
  - ▶ Scheduling methods decide which process to swap out


# Placement Algorithm with Partitions

- ▶ Unequal-size partitions: use of multiple queues
  - ▶ assign each process to the smallest partition within which it will fit
  - ▶ a queue for each partition size
  - ▶ tries to minimize internal fragmentation
  - ▶ Problem: some queues will be empty if no processes within a size range is present



# Placement Algorithm with Partitions

- ▶ Unequal-size partitions: use of a single queue
  - ▶ When its time to load a process into main memory the smallest available partition that will hold the process is selected
- ▶ increases the level of multiprogramming at the expense of internal fragmentation
  - ▶ The smallest available partition may be much greater than the process size



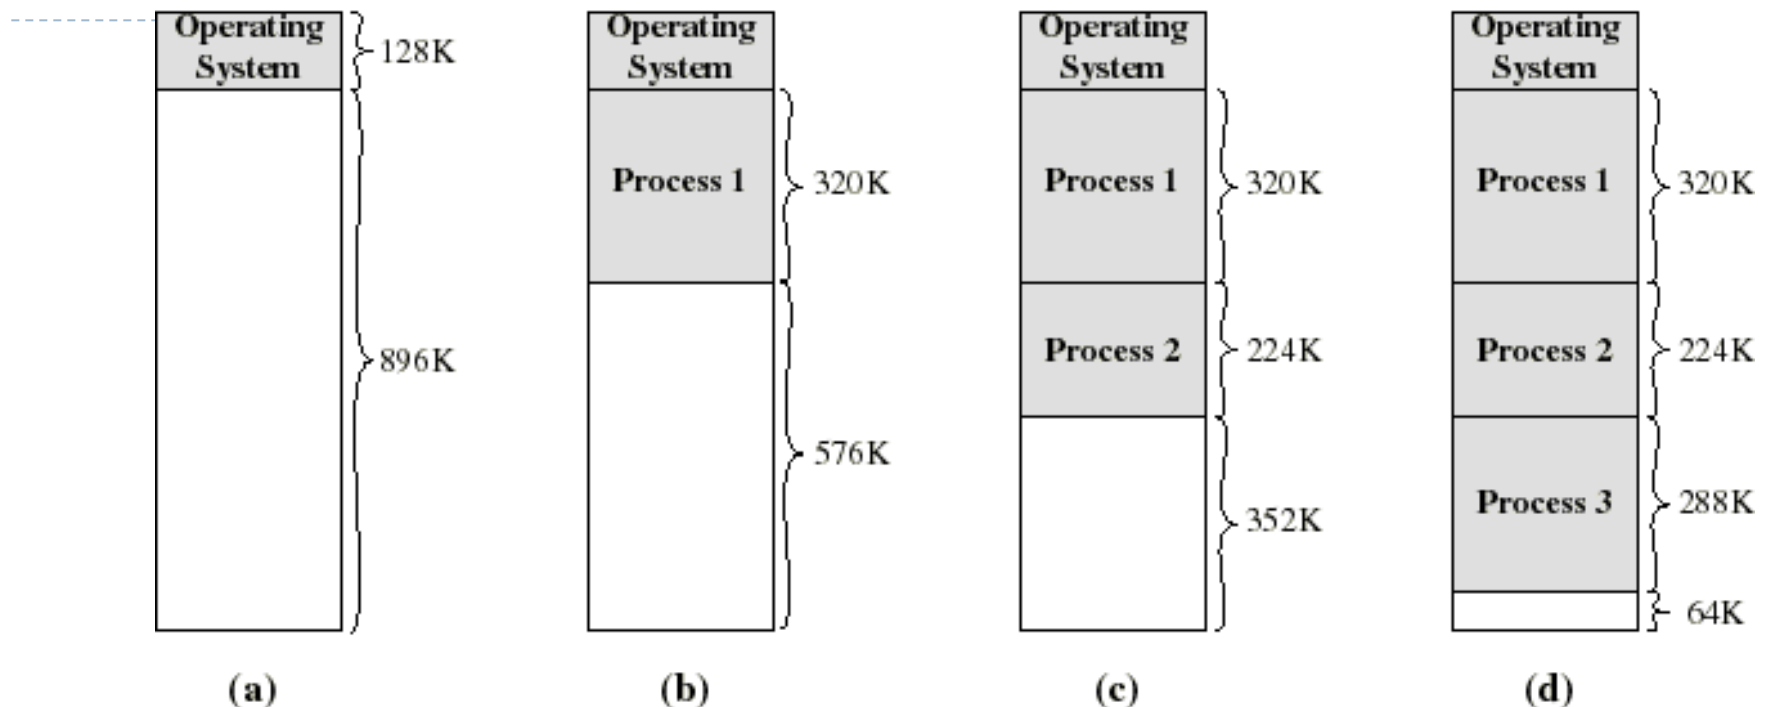
# Dynamic Partitioning

---

- ▶ Partitions are of variable length and number
- ▶ Each process is allocated exactly as much memory as it requires
- ▶ Eventually holes are formed in main memory. This is called **external fragmentation**
- ▶ Must use **compaction** to shift processes so they are contiguous and all free memory is in one block
- ▶ Used in IBM's OS/MVT (Multiprogramming with a Variable number of Tasks)

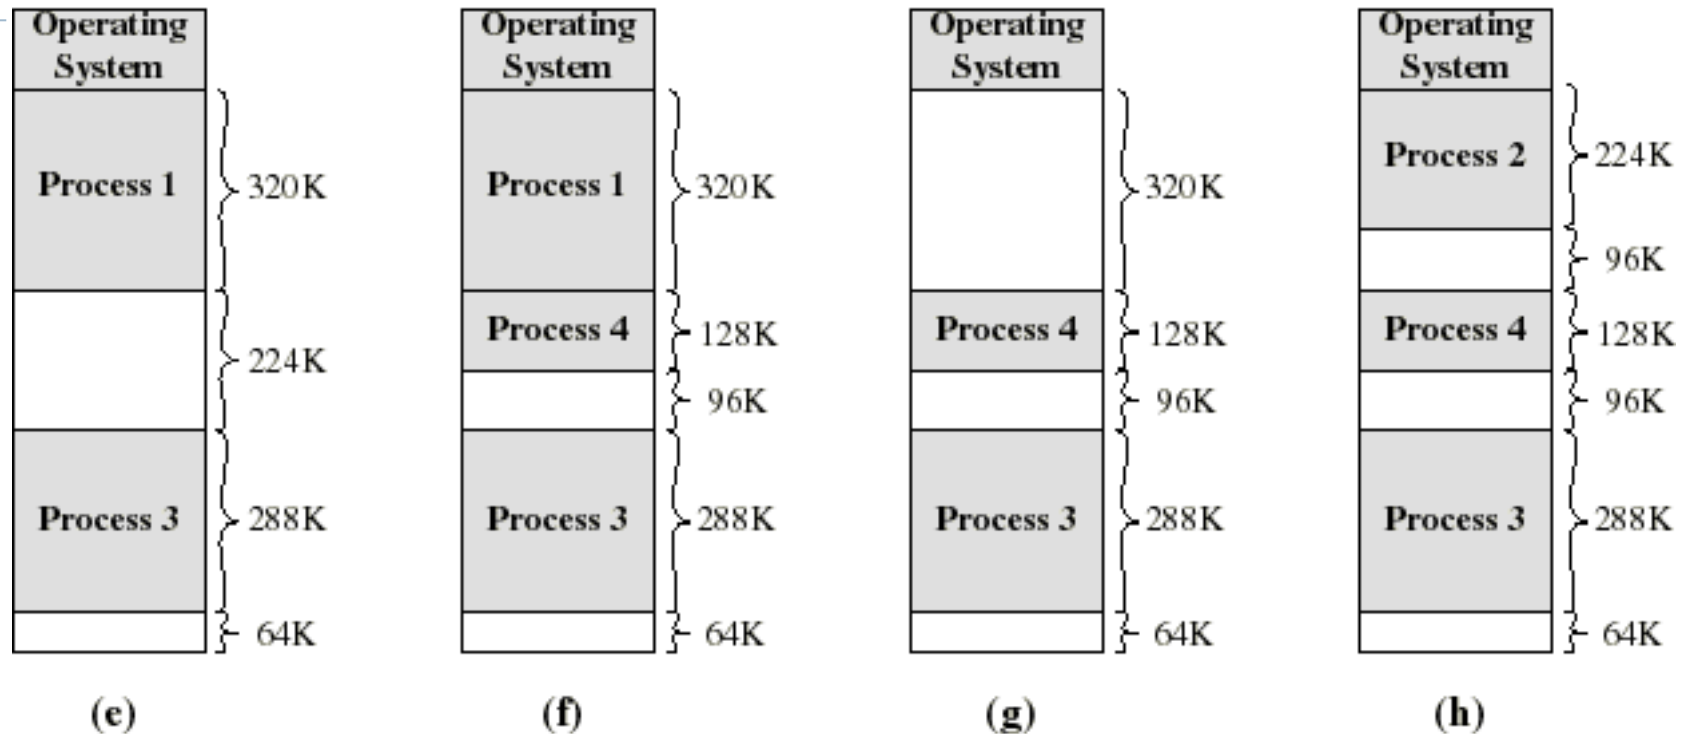


# Dynamic Partitioning: an example



- ▶ A hole of 64K is left after loading 3 processes: not enough room for another process
- ▶ Eventually each process is blocked. The OS swaps out process 2 to bring in process 4

# Dynamic Partitioning: an example



- ▶ another hole of 96K is created
- ▶ Eventually each process is blocked. The OS swaps out process 1 to bring in again process 2 and another hole of 96K is created...
- ▶ Compaction would produce a single hole of 256K

# Placement Algorithm

- ▶ Used to decide which free block to allocate to a process
- ▶ Goal: to reduce usage of compaction (time consuming)
- ▶ Possible algorithms:
  - ▶ **Best-fit**: choose smallest hole
  - ▶ **First-fit**: choose first hole from beginning
  - ▶ **Next-fit**: choose first hole from last placement




Figure 7.5 Example Memory Configuration before and after Allocation of 16-Mbyte Block

# Placement Algorithm: comments

---

- ▶ Next-fit often leads to allocation of the largest block at the end of memory
- ▶ First-fit favors allocation near the beginning: tends to create less fragmentation than Next-fit
- ▶ Best-fit searches for smallest block: the fragment left behind is small as possible
  - ▶ main memory quickly forms holes too small to hold any process: compaction generally needs to be done more often

# Buddy System

- ▶ Entire space available is treated as a single block of  $2^U$
- ▶ If a request of size  $s$  such that  $2^{U-1} < s \leq 2^U$ , entire block is allocated
  - ▶ Otherwise block is split into two equal buddies
  - ▶ Process continues until smallest block greater than or equal to  $s$  is generated

# Example of Buddy System

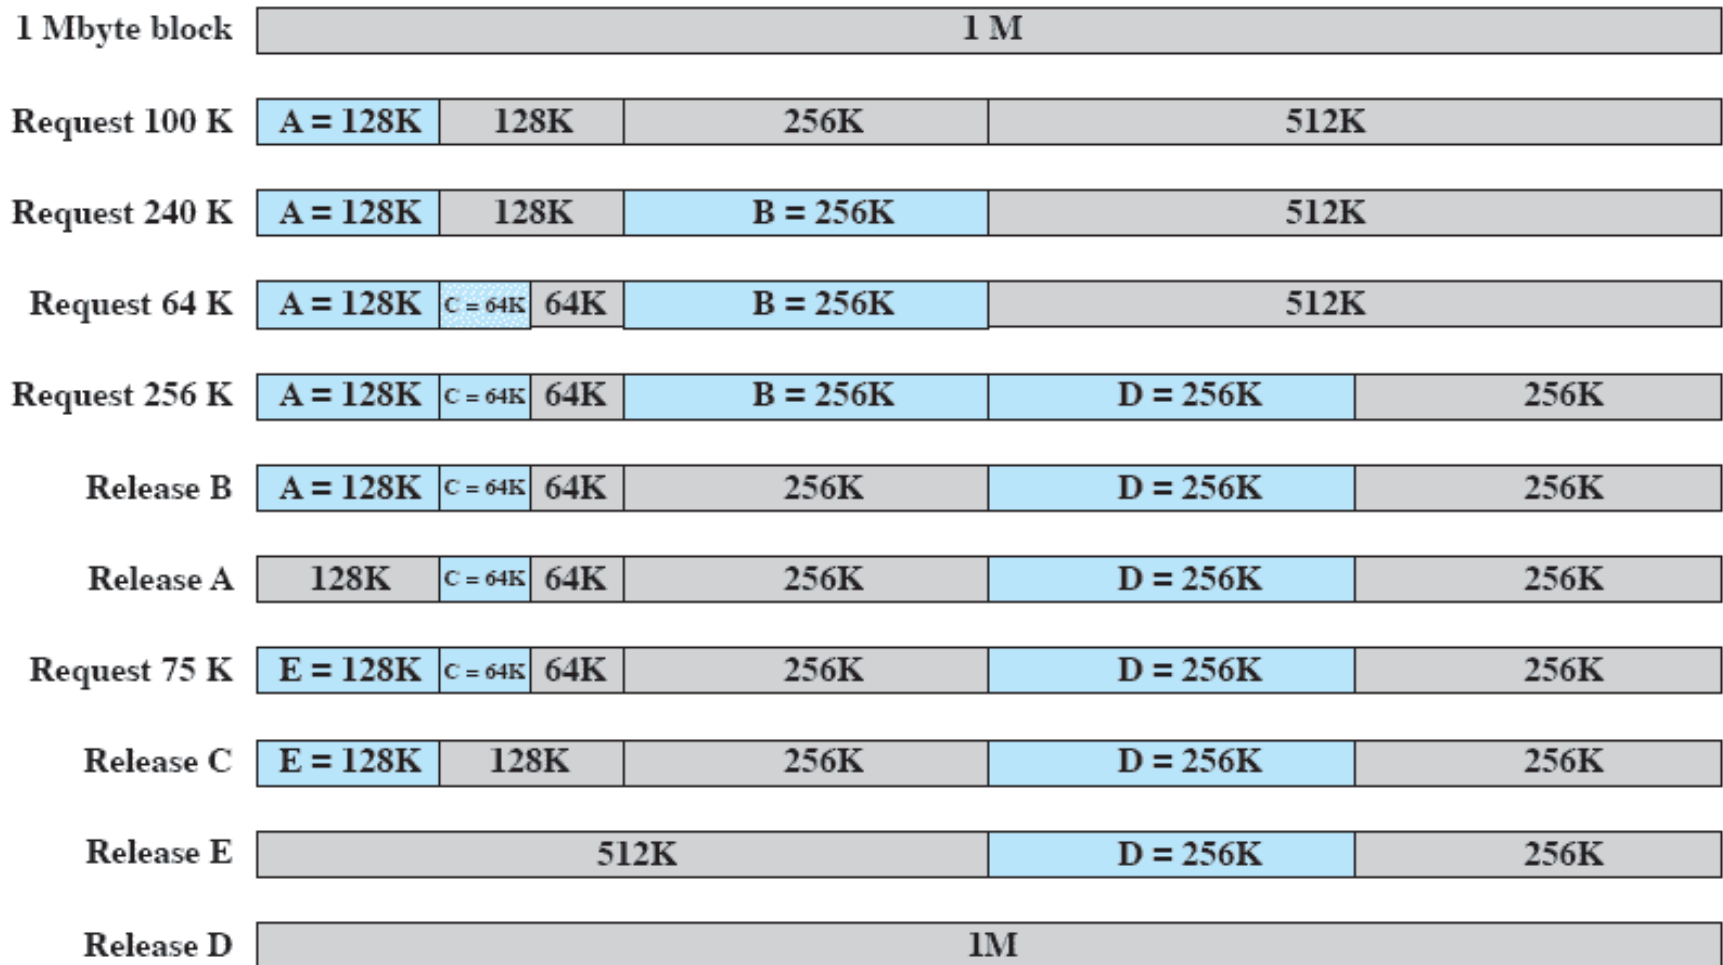


Figure 7.6 Example of Buddy System

# Tree Representation of Buddy System




Figure 7.7 Tree Representation of Buddy System

# Replacement Algorithm

---

- ▶ When all processes in main memory are blocked, the OS must choose which process to replace
  - ▶ A process must be swapped out (to a Blocked-Suspend state) and be replaced by a new process or a process from the Ready-Suspend queue
  - ▶ We will discuss later such algorithms for memory management schemes using virtual memory



# Relocation

---

- ▶ When program is loaded into memory, the actual (absolute) memory locations are determined
- ▶ Because of swapping and compaction, a process may occupy different main memory locations during its lifetime
- ▶ Hence physical memory references by a process cannot be fixed
- ▶ This problem is solved by distinguishing between **logical address** and **physical address**

# Address Types

---

- ▶ A **physical address** (absolute address) is a physical location in main memory
- ▶ A **logical address** is a reference to a memory location independent of the physical structure/organization of memory
- ▶ Compilers produce code in which all memory references are logical addresses
- ▶ A **relative address** is an example of logical address in which the address is expressed as a location relative to some known point in the program (ex: the beginning)

# Address Translation

---

- ▶ Relative address is the most frequent type of logical address used in pgm modules
- ▶ Such modules are loaded in main memory with all memory references in relative form
- ▶ Physical addresses are calculated “on the fly” as the instructions are executed
- ▶ This is called dynamic run-time loading
- ▶ For adequate performance, the translation from relative to physical address must be done by hardware

# Hardware translation of addresses

---

- ▶ When a process is assigned to the running state, a base register (in CPU) gets loaded with the starting physical address of the process
- ▶ A bound register gets loaded with the process's ending physical address
- ▶ When a relative addresses is encountered, it is added with the content of the base register to obtain the physical address which is compared with the content of the bound register
- ▶ This provides hardware protection: each process can only access memory within its process image

# Example Hardware for Address Translation

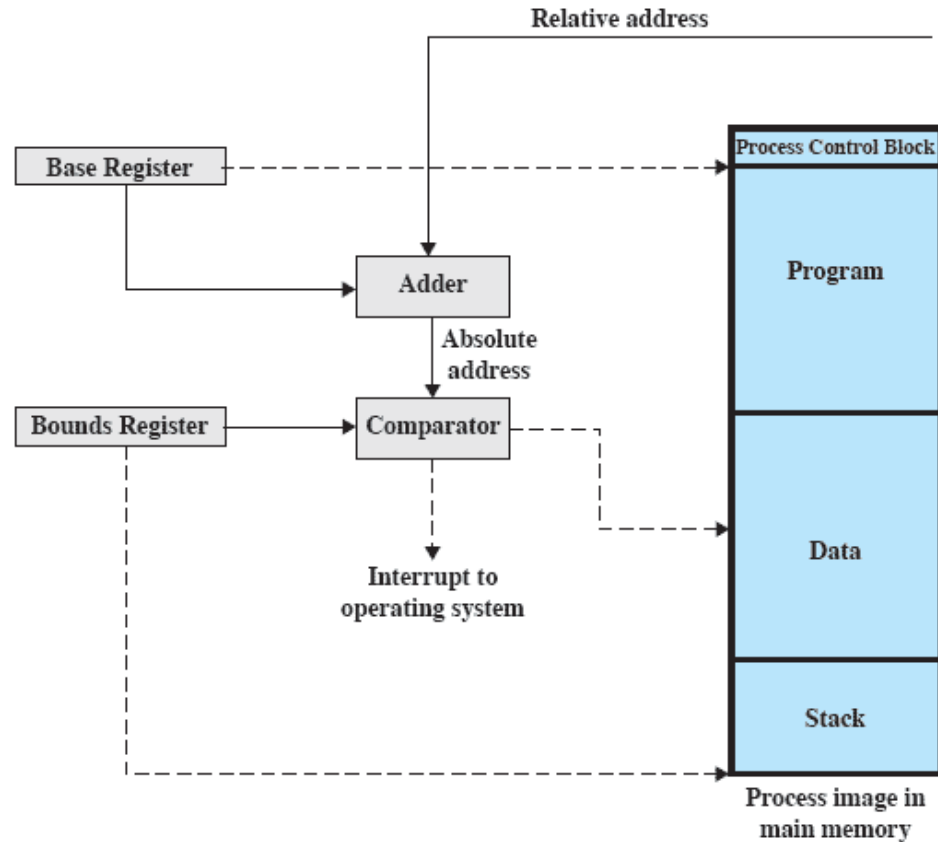


Figure 7.8 Hardware Support for Relocation

# Paging

- ▶ Partition memory into small equal fixed-size chunks and divide each process into the same size chunks
- ▶ The chunks of a process are called pages and chunks of memory are called frames
- ▶ Consequence: a process does not need to occupy a contiguous portion of memory

# Paging

- ▶ Operating system maintains a page table for each process
  - ▶ Contains the frame location for each page in the process
  - ▶ Memory address consists of a page number and offset within the page

# Example of process loading

Frame number	Main memory
0	
1	
2	
3	
4	
5	
6	
7	
8	
9	
10	
11	
12	
13	
14	

(a) Fifteen Available Pages

Frame number	Main memory
0	A.0
1	A.1
2	A.2
3	A.3
4	
5	
6	
7	
8	
9	
10	
11	
12	
13	
14	

(b) Load Process A

Frame number	Main memory
0	A.0
1	A.1
2	A.2
3	A.3
4	B.0
5	B.1
6	B.2
7	
8	
9	
10	
11	
12	
13	
14	

(b) Load Process B

Frame number	Main memory
0	A.0
1	A.1
2	A.2
3	A.3
4	B.0
5	B.1
6	B.2
7	C.0
8	C.1
9	C.2
10	C.3
11	
12	
13	
14	

(d) Load Process C

- Now suppose that process B is swapped out



# Example of process loading (cont.)

- ▶ When process A and C are blocked, the pager loads a new process D consisting of 5 pages
- ▶ Process D does not occupied a contiguous portion of memory
- ▶ There is no external fragmentation
- ▶ Internal fragmentation consist only of the last page of each process

Main memory	
0	A.0
1	A.1
2	A.2
3	A.3
4	
5	
6	
7	C.0
8	C.1
9	C.2
10	C.3
11	
12	
13	
14	

(e) Swap out B

Main memory	
0	A.0
1	A.1
2	A.2
3	A.3
4	D.0
5	D.1
6	D.2
7	C.0
8	C.1
9	C.2
10	C.3
11	D.3
12	D.4
13	
14	

(f) Load Process D

# Page Tables

0	0
1	1
2	2
3	3

**Process A**  
page table

0	—
1	—
2	—

**Process B**  
page table

0	7
1	8
2	9
3	10

**Process C**  
page table

0	4
1	5
2	6
3	11
4	12

**Process D**  
page table

13
14

**Free frame**  
list

- ▶ The OS now needs to maintain (in main memory) a **page table** for each process
- ▶ Each entry of a page table consist of the frame number where the corresponding page is physically located
- ▶ The page table is indexed by the page number to obtain the frame number
- ▶ A free frame list, available for pages, is maintained

# Logical address used in paging

---


- ▶ Within each program, each logical address must consist of a **page number** and an **offset** within the page
- ▶ A CPU register always holds the starting physical address of the page table of the currently running process
- ▶ Presented with the logical address (page number, offset) the processor accesses the page table to obtain the physical address (frame number, offset)

# Logical address in paging

- ▶ The logical address becomes a relative address when the page size is a power of 2
- ▶ Ex: if 16 bits addresses are used and page size = 1K, we need 10 bits for offset and have 6 bits available for page number
- ▶ Then the 16 bit address obtained with the 10 least significant bit as offset and 6 most significant bit as page number is a location relative to the beginning of the process

Logical address =  
Page# = 1, Offset = 478

0000010111011110

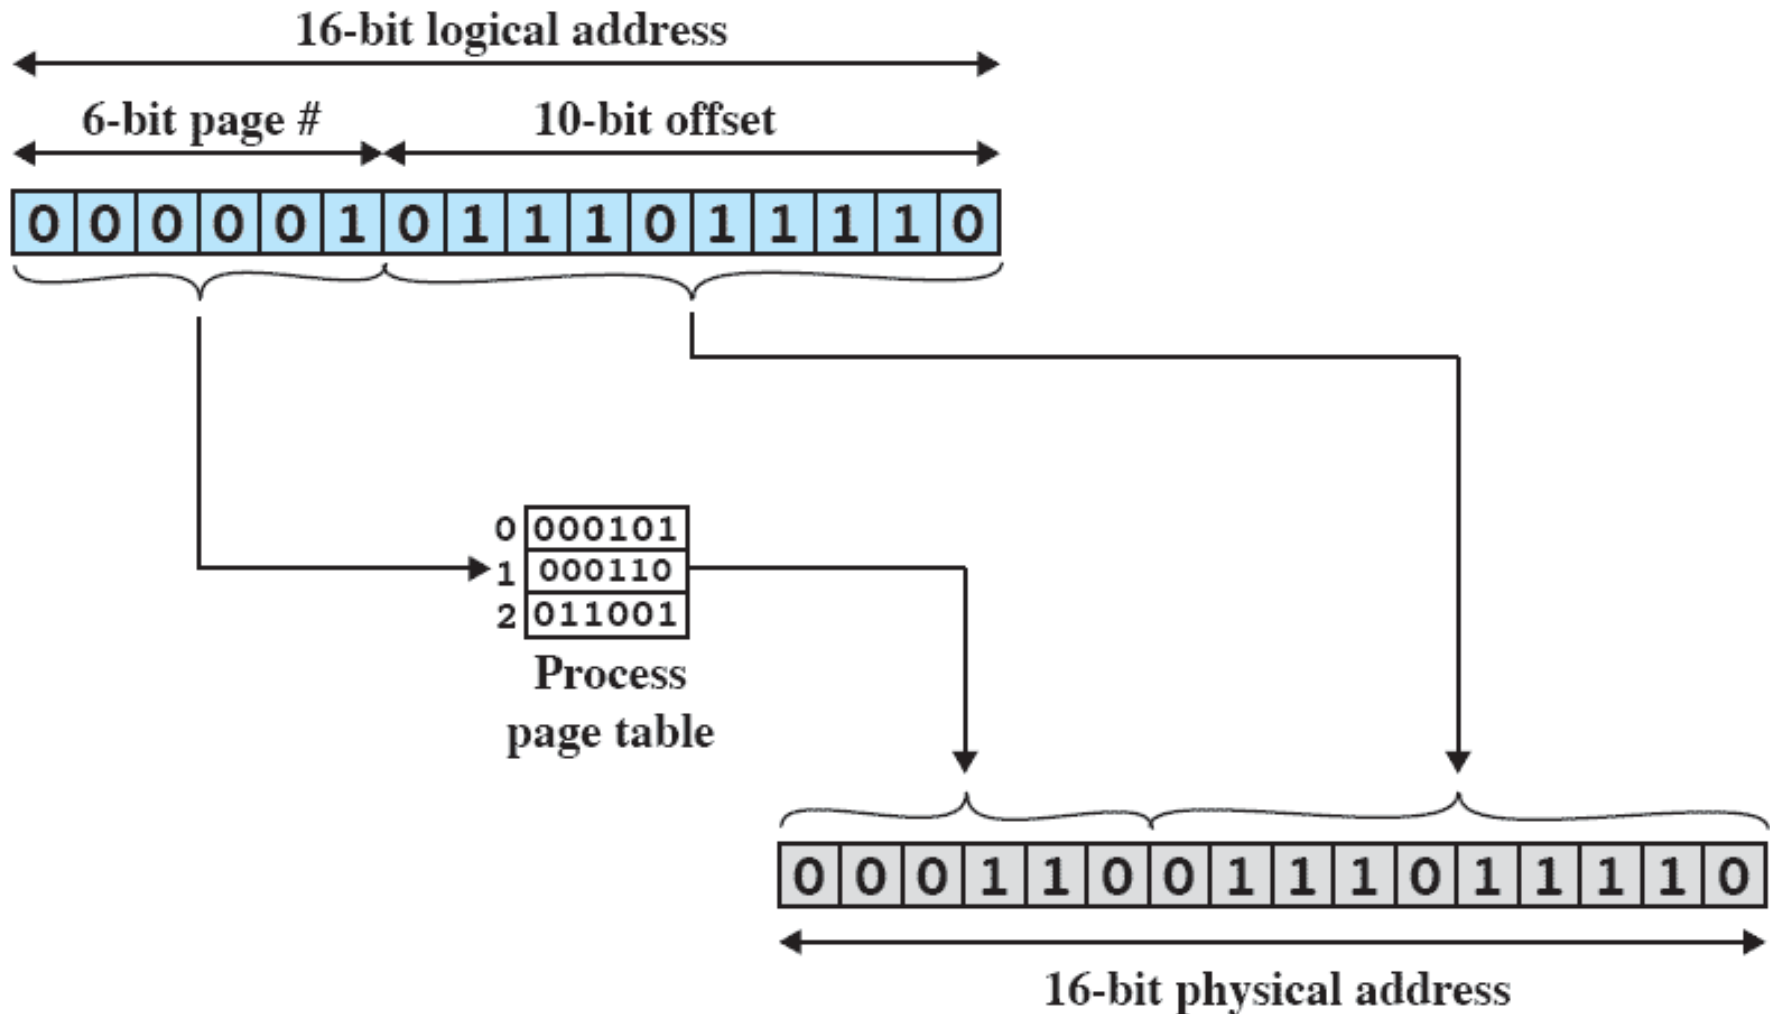


# Logical address in paging

---

- ▶ By using a page size of a power of 2, the pages are invisible to the programmer, compiler/assembler, and the linker
- ▶ Address translation at run-time is then easy to implement in hardware
  - ▶ logical address  $(n,m)$  gets translated to physical address  $(k,m)$  by indexing the page table and appending the same offset  $m$  to the frame number  $k$

# Logical-to-Physical Address Translation in Paging



(a) Paging

# Simple Segmentation

---

- ▶ Each program is subdivided into blocks of non-equal size called **segments**
- ▶ There is a maximum segment length
- ▶ When a process gets loaded into main memory, its different segments can be located anywhere
- ▶ Each segment is fully packed with instructions/data: no internal fragmentation
- ▶ There is external fragmentation; it is reduced when using small segments

# Simple Segmentation

---

- ▶ In contrast with paging, segmentation is visible to the programmer
  - ▶ provided as a convenience to organize logically programs (ex: data in one segment, code in another segment)
  - ▶ must be aware of segment size limit
- ▶ The OS maintains a **segment table** for each process. Each entry contains:
  - ▶ the starting physical addresses of that segment.
  - ▶ the length of that segment (for protection)




# Logical address used in segmentation

---

- ▶ When a process enters the Running state, a CPU register gets loaded with the starting address of the process's segment table.
- ▶ Presented with a logical address (segment number, offset) =  $(n, m)$ , the CPU indexes (with  $n$ ) the segment table to obtain the starting physical address  $k$  and the length  $l$  of that segment
- ▶ The physical address is obtained by adding  $m$  to  $k$  (in contrast with paging)
  - ▶ the hardware also compares the offset  $m$  with the length  $l$  of that segment to determine if the address is valid
  - ▶ we cannot directly obtain the logical address from the physical address (in contrast with paging)

# Logical-to-Physical Address Translation in segmentation



# Simple segmentation and paging comparison

---

- ▶ Segmentation requires more complicated hardware for address translation
- ▶ Segmentation suffers from external fragmentation
- ▶ Paging only yield a small internal fragmentation
- ▶ Segmentation is visible to the programmer whereas paging is transparent
- ▶ Segmentation can be viewed as commodity offered to the programmer/compiler to organize logically a program into segments and using different kinds of protection (ex: execute-only for code but read-write for data)
  - ▶ for this we need to use protection bits in segment table entries