

## Chapter 8 : Main Memory



## Background

- Program brought (from disk) into memory and placed within a process for it to be run
- Main memory and registers are only storage CPU can access directly
- Register access in one CPU clock (or less)
- Main memory can take many cycles
- Cache sits between main memory and CPU registers
- Protection of memory required to ensure correct operation


Operating System Concepts – 7th Edition, Feb 22, 2005

8.2

Silberschatz, Galvin and Gagne ©2005

## Base and Limit Registers

- A pair of **base** and **limit** registers define the logical address space



Operating System Concepts – 7th Edition, Feb 22, 2005

8.3

Silberschatz, Galvin and Gagne ©2005

## Binding of Instructions and Data to Memory


- Address binding (instructions/data) to memory addresses can happen at:
  - **Compile time**: memory location known a priori, **absolute code** generated; must recompile if starting location changes
  - **Load time**: generate **relocatable code** if memory location not known at compile time
  - **Execution time**: **binding** delayed until **run time**. Process can be moved during execution from one memory segment to another. Need **hardware support** for address maps (e.g., base and limit registers)

Operating System Concepts – 7th Edition, Feb 22, 2005

8.4

Silberschatz, Galvin and Gagne ©2005

## Multistep Processing of a User Program



Operating System Concepts – 7th Edition, Feb 22, 2005

8.5

Silberschatz, Galvin and Gagne ©2005

## Logical vs. Physical Address Space

- Concept: **logical address space** bound to separate **physical address space**
  - **Logical address** – generated by the CPU (also referred to as **virtual address**)
  - **Physical address** – address seen by the memory unit
- Logical and physical addresses:
  - **same** in compile-time/load-time address-binding schemes;
  - **different** in run-time address-binding scheme

Operating System Concepts – 7th Edition, Feb 22, 2005

8.6


Silberschatz, Galvin and Gagne ©2005

## Memory-Management Unit (MMU)

- **Hardware** device that maps virtual to physical address
- Value in **relocation register** added to every address generated by a user process before sending to memory
- User program deals with **logical** addresses; it never sees the *real* physical addresses

7

## Dynamic relocation using a relocation register



8

## Dynamic Loading

- Routine is not loaded until it is **called**
- Better **memory-space utilization**; unused routine is never loaded
- Useful when **large** amounts of code are needed **infrequently**
- No special support from the operating system is required; implemented through **program design**

9

## Dynamic Linking

- **Linking** postponed until **run-time**
- Small piece of code, **stub**, **locate** the appropriate library routine (**memory-resident**)
- Stub **replaces itself** with the address of the routine, and executes the routine
- OS: **check** if routine is in processes' memory address
- Dynamic linking is particularly useful for libraries
- System also known as **shared libraries**


10

## Swapping

- Process swapped temporarily from memory to **backing store**, then back into memory for continued execution
- **Fast disk**; large to accommodate copies of **all memory images** for all users; direct access to these memory images
- **Roll out, roll in** –variant used for **priority-based** scheduling (low-priority process swapped out; high-priority executed)
- **Transfer time**: directly proportional to amount of memory swapped
- Modified versions found on many systems (UNIX, Linux, Windows)
- System maintains a **ready queue** (swapped) of processes with memory images on disk

11

## Schematic View of Swapping



12

## Contiguous Allocation

- Main memory divided into **two partitions**:
  - Resident **OS**, held in **low memory** (with interrupt vector)
  - User **processes** then held in **high memory**
- Relocation registers** used to **protect** user processes from each other (and also OS code and data)
  - Base register** contains value of **smallest physical address**
  - Limit register** contains **range of logical addresses** – each logical address must be less than the limit register
  - MMU maps logical address *dynamically*


Operating System Concepts – 7th Edition, Feb 22, 2005

8.13

Silberschatz, Galvin and Gagne ©2005

13

## HW address protection with base and limit registers



Operating System Concepts – 7th Edition, Feb 22, 2005


8.14

Silberschatz, Galvin and Gagne ©2005

14

## Contiguous Allocation (Cont.)

- Multiple-partition allocation
  - Hole** – block of available memory; holes of various size scattered throughout memory
- New process: memory **allocated from a hole** large enough to accommodate it
- OS maintains information about:
  - a) **allocated partitions**
  - b) **free partitions (hole)**



Operating System Concepts – 7th Edition, Feb 22, 2005

8.15

Silberschatz, Galvin and Gagne ©2005

15

## Dynamic Storage-Allocation Problem

How to satisfy a request of size  $n$  from a list of free holes?

- First-fit**: Allocate the *first* hole that is big enough
- Best-fit**: Allocate the *smallest* hole that is big enough; must search entire list, unless ordered by size
  - Produces the *smallest* leftover hole
- Worst-fit**: Allocate the *largest* hole; must also search entire list
  - Produces the *largest* leftover hole

First-fit and best-fit better than worst-fit in terms of speed and storage utilization

Operating System Concepts – 7th Edition, Feb 22, 2005

8.16

Silberschatz, Galvin and Gagne ©2005

16

## Fragmentation

- External Fragmentation** – total memory space exists to satisfy a request, but *not contiguous*
- Internal Fragmentation** – allocated memory slightly larger than requested memory; size difference is memory *internal to a partition*, but not being used
- Reduce external fragmentation by **compaction**
  - Shuffle memory contents; place free memory together in *one large block*
  - Compaction possible *only* if relocation is dynamic; done at runtime
  - I/O problem
    - Latch job in memory while it is involved in I/O
    - Do I/O only into OS buffers

Operating System Concepts – 7th Edition, Feb 22, 2005

8.17

Silberschatz, Galvin and Gagne ©2005

17

## Paging

- Logical** address space of a process can be **noncontiguous**; process allocated physical memory whenever available
- Divide physical memory into **fixed-sized** blocks called **frames** (size is power of 2, between 512 bytes and 8,192 bytes)
- Divide logical memory into blocks of same size called **pages**
- Keep track of all free frames
- To run a program of  $n$  pages, find  $n$  free frames and load program
- Set up a **page table** to translate logical to physical addresses
- Internal fragmentation**

Operating System Concepts – 7th Edition, Feb 22, 2005

8.18


Silberschatz, Galvin and Gagne ©2005

18

## Address Translation Scheme

- Address generated by CPU is divided into:
  - Page number ( $p$ )** – index into *page table*.  
Contains base address of *each page* in physical memory

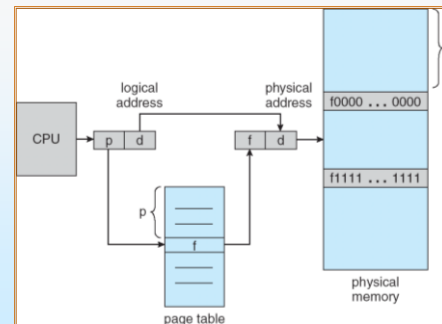
- Page offset ( $d$ )** – combined with base address to define physical memory address sent to the memory unit



- For given logical address space  $2^m$  and page size  $2^n$

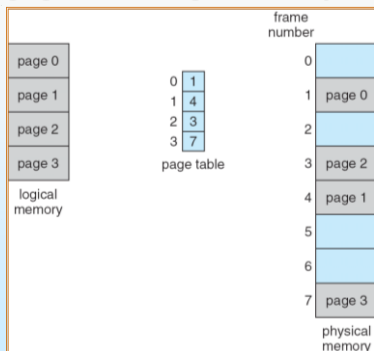
19

## Paging Hardware




20

## Paging Model of Logical and Physical Memory



21


## Paging Example



32-byte memory and 4-byte pages

22

## Free Frames



Before allocation

After allocation

23

## Implementation of Page Table

- Page table is kept in *main memory*
- Page-table base register (PTBR) points to the **page table**
- Page-table length register (PRLR) indicates **size** of the page table
- In this scheme every data/instruction access: two memory accesses.  
One for the page table and one for the data/instruction.
- Two memory access problem can be solved using a special fast-lookup hardware cache called **associative memory** or **translation look-aside buffers (TLBs)**
- Some TLBs store **address-space identifiers (ASIDs)** in each TLB entry – uniquely identifies each process to provide address-space protection for that process

24

## Associative Memory

- Associative memory – parallel search


Page #	Frame #

Address translation (p, d)

- If p is in associative register, get frame # out
- Otherwise get frame # from page table in memory

25

## Paging Hardware With TLB



26

## Effective Access Time

- Associative Lookup =  $\epsilon$  time unit
- Assume memory cycle time is 1 microsecond
- Hit ratio – percentage of times that a page number found in associative registers; ratio related to number of associative registers
- Hit ratio =  $\alpha$

- Effective Access Time (EAT)

$$\begin{aligned} \text{EAT} &= (1 + \epsilon) \alpha + (2 + \epsilon)(1 - \alpha) \\ &= 2 + \epsilon - \alpha \end{aligned}$$

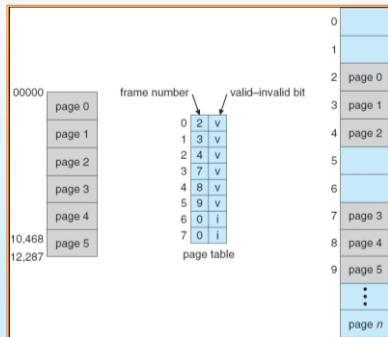
27

## Memory Protection

- Implemented by associating protection bit with each frame
- Valid-invalid** bit attached to each entry in the page table:
  - "valid": associated page is in the process' logical address space (legal page)
  - "invalid": page is not in the process' logical address space

28

## Valid (v) or Invalid (i) Bit In A Page Table

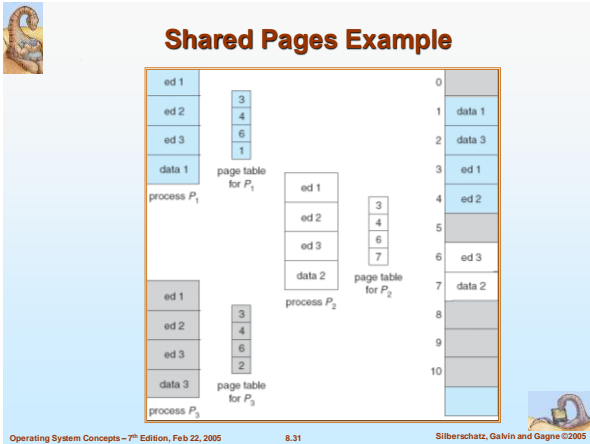


29

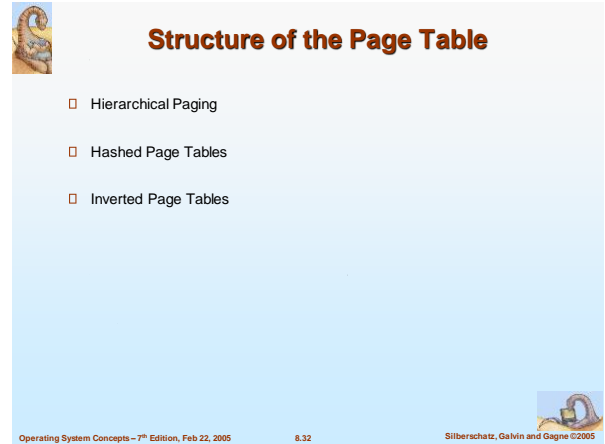
## Shared Pages

- Shared code**
  - One copy of read-only (**reentrant**) code shared among processes (i.e., text editors, compilers, window systems).
  - Shared code must appear in *same location* in the *logical address space* of all processes
- Private code and data**
  - Each process keeps a *separate copy* of the code and data
  - Pages for the private code and data can appear *anywhere* in the logical address space

30



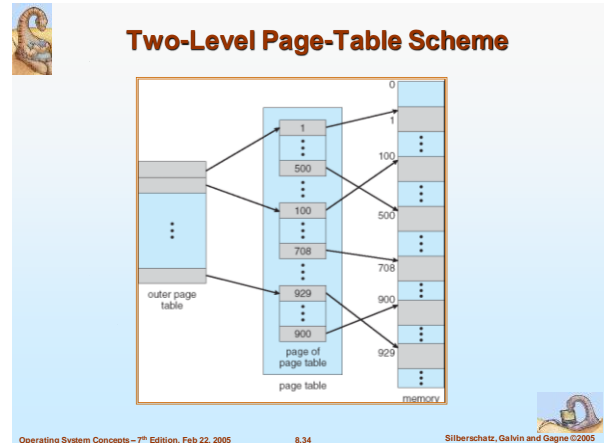
31



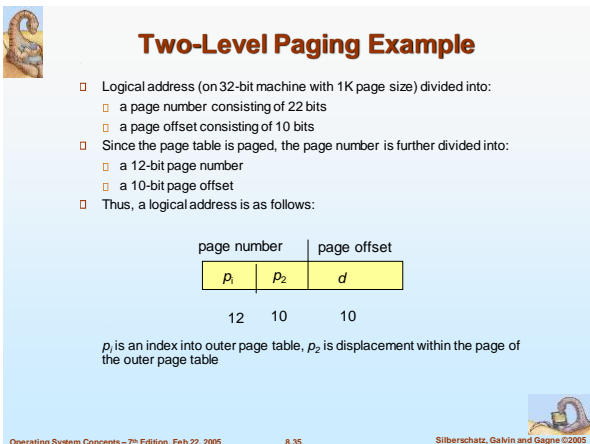
32



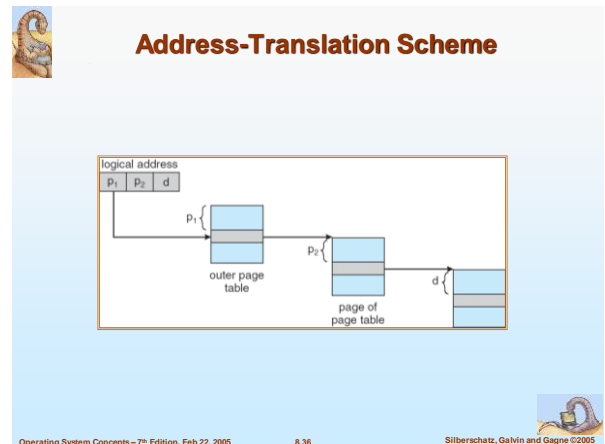
33



34



35



36

## Segmentation

- Memory-management scheme that supports user view of memory
- Program: collection of segments. Segment: a logical unit such as:
  - main program,
  - procedure,
  - function,
  - method,
  - object,
  - local variables, global variables,
  - common block,
  - stack,
  - symbol table, arrays


Operating System Concepts – 7th Edition, Feb 22, 2005

8.42

Silberschatz, Galvin and Gagne ©2005

42

## User's View of a Program




Operating System Concepts – 7th Edition, Feb 22, 2005

8.43

Silberschatz, Galvin and Gagne ©2005

43

## Logical View of Segmentation



Operating System Concepts – 7th Edition, Feb 22, 2005

8.44

Silberschatz, Galvin and Gagne ©2005

44

## Segmentation Architecture

- Logical address consists of a two tuple:
  - <segment-number, offset>
- **Segment table** – maps two-dimensional physical addresses; each table entry has:
  - **base** – starting physical address where segments reside in memory
  - **limit** – specifies the length of the segment
- **Segment-table base register (STBR)**: points to segment table's location in memory
- **Segment-table length register (STLR)**: indicates number of segments used by a program;
  - segment number  $s$  is legal if  $s < \text{STLR}$

Operating System Concepts – 7th Edition, Feb 22, 2005

8.45

Silberschatz, Galvin and Gagne ©2005

45

## Segmentation Architecture (Cont.)

- Protection
  - With each entry in segment table associate:
    - ▶ validation bit = 0  $\Rightarrow$  illegal segment
    - ▶ read/write/execute privileges
- Protection bits associated with segments; code sharing occurs at segment level
- Since segments vary in length, memory allocation is a dynamic storage-allocation problem
- A segmentation example is shown in the following diagram

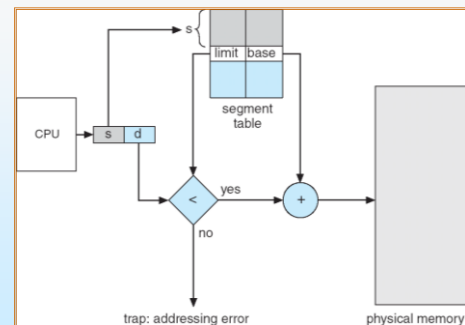
Operating System Concepts – 7th Edition, Feb 22, 2005

8.46

Silberschatz, Galvin and Gagne ©2005

46

## Segmentation Hardware

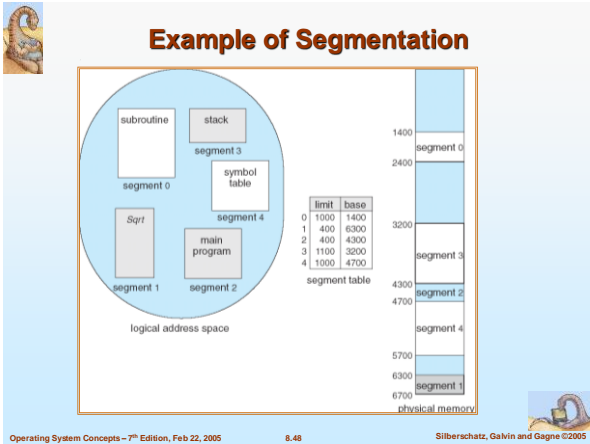


Operating System Concepts – 7th Edition, Feb 22, 2005

8.47

Silberschatz, Galvin and Gagne ©2005

47



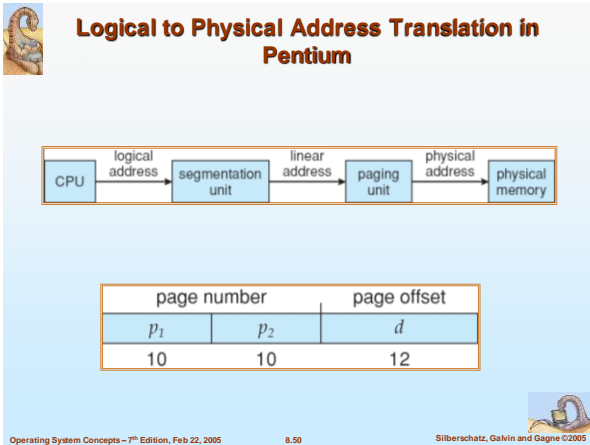
48

### Example: The Intel Pentium

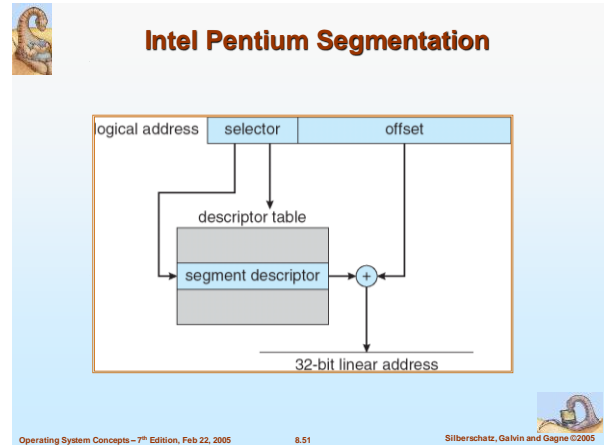
- Supports both segmentation and segmentation with paging
- CPU generates logical address
  - Given to segmentation unit
    - Which produces linear addresses
  - Linear address given to paging unit
    - Which generates physical address in main memory
  - Paging units form equivalent of MMU

Operating System Concepts – 7th Edition, Feb 22, 2005 8.49 Silberschatz, Galvin and Gagne ©2005


49



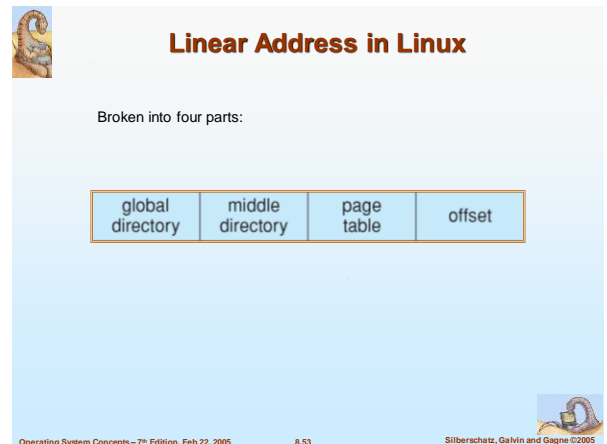
50



51



52




53

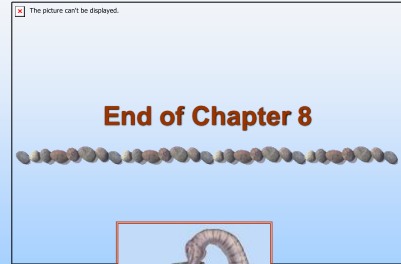




## Three-level Paging in Linux



54



55